

# **Autonomous Algorithmic Collusion**

Efficient Learning with Q-Function Approximation

**Lindsay Robinson**

*University of Melbourne*

# Motivation

Autonomous Algorithmic Collusion

## 'Autonomous Algorithmic Collusion'

- Firms delegate pricing to reinforcement learning algorithms
- Algorithms learn to tacitly collude through repeated interaction

## Reinforcement Learning Algorithms:

- Learn dynamic best response to non-learning competitors
- Strategically rationalizable if:
  1. Learn efficiently enough that losses from suboptimal play during learning period are negligible
  2. Learn to collude with learning competitors

	Non-Learning	Learning
Non-Learning	$\pi, \pi$	$\pi, \pi_{BR}$
Learning	$\pi_{BR}, \pi$	??

## Existing Evidence:

- Tabular Q-learning algorithms
  - Calvano et al. (2020 AER; 2023 IJIO)
  - Klein (2021 RJE)
  - Asker, Fershtman & Pakes (2022 AEAPP; 2023 JEMS)
  - Johnson, Rhodes & Wildenbeest (2023 E)
  - Abada & Lambin (2023 MS)
- Deep Q-learning algorithms
  - Hettich (2021)
  - Schlechtinger et al. (2024)
  - Dawid, Harting & Neugart (2024)
  - Deng, Schiffer & Bichler (2025)

## This Paper:

- Linear Q-learning algorithms
  - More efficient learning  
     $\implies$  Strategically rationalizable  
     $\implies$  Policy relevant

# Motivation

## Q-Learning Pricing Algorithms

### Repeated Pricing Games

- For firm  $f$ , at time period  $t - 1$ :
  - State = Price history:  $\mathbf{P}_{1:t-1}$
  - Action = Own price:  $p_{ft}$
  - Reward = Own profit:  $\pi_{ft}$

### Q-Learning Pricing Algorithms

- Maintain estimate of conditional value function - the ‘Q-function’:  $Q(p_{ft} | \mathbf{P}_{1:t-1})$

- Recursive Bellman equation:

$$Q(p_{ft} | \mathbf{P}_{1:t-1}) := \underbrace{\pi_{ft}(p_{ft})}_{\text{period profit}} + \delta \underbrace{\max_{P_{ft+1}} Q(p_{ft+1} | \mathbf{P}_{1:t}(p_{ft}))}_{\text{continuation profits}}$$

- Each period, observe  $(\mathbf{P}_{1:t-1}, p_{ft}, \pi_{ft}, \mathbf{P}_{1:t})$  and update estimated  $\widehat{Q}$ :

- Gradient step in direction that minimises Bellman error:

$$\widehat{e}_b = \underbrace{(\pi_{ft} + \delta \max_{P_{ft+1}} \widehat{Q}(p_{ft+1} | \mathbf{P}_{1:t}))}_{\text{realisation}} - \underbrace{\widehat{Q}(p_{ft} | \mathbf{P}_{1:t-1})}_{\text{estimate}}$$

# Motivation

## Q-Learning Pricing Algorithms - Q-Function Representations

### Tabular

- Slow, inefficient learning
  - No generalisation from similar experience
- Curse of dimensionality

$$Q(p_{ft} | P_{1:t-1}) = \theta_{p_{ft}} p_{1:t-1}$$

$$\Theta = \begin{bmatrix} \theta_{11} & \dots & \theta_{1|\mathcal{P}|} \\ \vdots & \ddots & \vdots \\ \theta_{|\mathcal{P}|1} & \dots & \theta_{|\mathcal{P}||\mathcal{P}|} \end{bmatrix}$$

### Deep

- Delayed, outdated learning
  - Backpropagation stability requires experience bank
  - Optimize against old competitor strategies

$$Q(p_{ft} | P_{1:t-1}) = f(\omega, f(\omega, \dots f(\omega, \phi(p_{ft}, P_{1:t-1}))))$$

### Linear

- Efficient learning
  - Generalisation from similar experience
  - Parsimonious parametrization
- Linear update rule
  - Immediate learning, reactivity

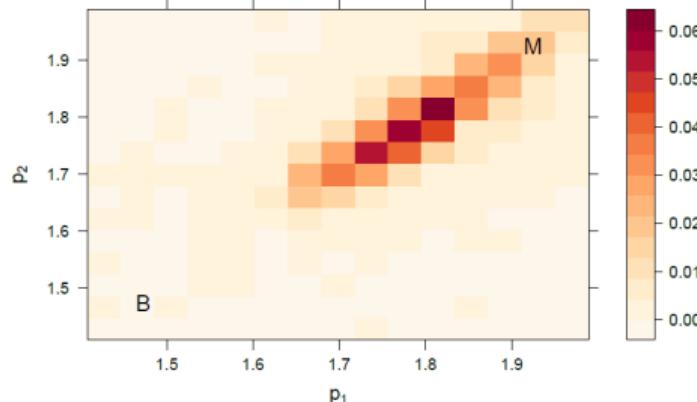
$$Q(p_{ft} | P_{1:t-1}) = \omega' \cdot \phi(p_{ft}, P_{1:t-1}) = \begin{bmatrix} \omega_1 \\ \vdots \\ \omega_W \end{bmatrix}' \cdot \begin{bmatrix} \phi_1(p_{ft}, P_{1:t-1}) \\ \vdots \\ \phi_W(p_{ft}, P_{1:t-1}) \end{bmatrix}$$

# Motivation

Q-Learning Pricing Algorithms - Q-Function Representations - Practical Significance

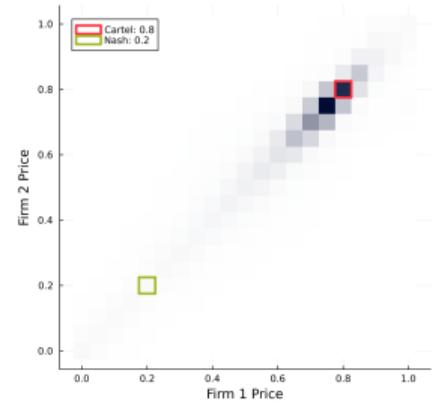
## Calvano et al (2020 AER):

- **Q-Function:** Tabular
  - Less efficient learning
- **State:** Most recent prices
  - Myopic conditioning/information set
  - Limited strategy space
- **Initialisation:** Pre-trained offline
  - Requires ex-ante demand, competitor knowledge
- **Periods before Collusion:** ‘Millions’
  - Extended algorithm commitment



## Current Paper:

- **Q-Function:** Linear
  - More efficient learning
- **State:** Price history
  - Full information set
  - Richer strategy space
- **Initialisation:** Naive
  - No required demand, competitor knowledge
- **Periods before Collusion:**  $\sim 3000$ 
  - Shorter algorithm commitment



# Setup

## Strategic Environment

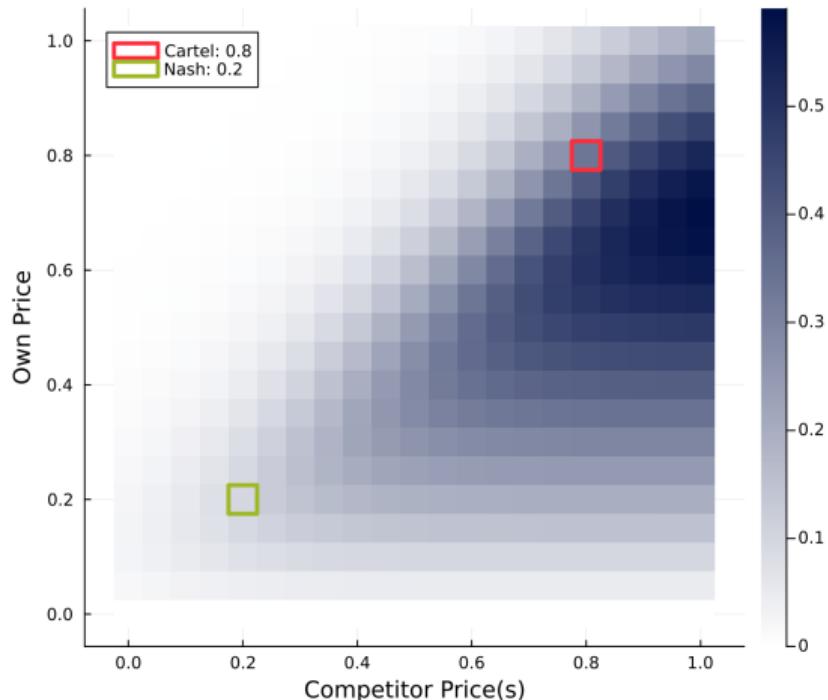
### Strategic Environment:

- Firms:  $f = 1, 2, \dots, F$
- Time periods:  $t = 1, 2, \dots, T$
- Action space:  $p \in \mathcal{P} = \{0, \dots, 1\}$ 
  - $|\mathcal{P}| = A$
- Marginal costs:  $c_f = 0 \forall f$
- Logit demand:

$$\pi_f(p_f, p_{f^-}) = p_f \cdot \frac{\exp\left(\frac{1-p_f}{\mu}\right)}{\exp\left(\frac{\eta}{\mu}\right) + \sum_{f=1}^F \exp\left(\frac{1-p_f}{\mu}\right)}$$

- $\mu$ : Market power
  - Product differentiation
- $\eta$ : Aggregate demand elasticity
  - Outside good strength

Own Profit  
( $F = 2, A = 21, \eta = 0.1, \mu = 0.1$ )



# Setup

## Q-Function Approximation

- Linear Q-function with features:

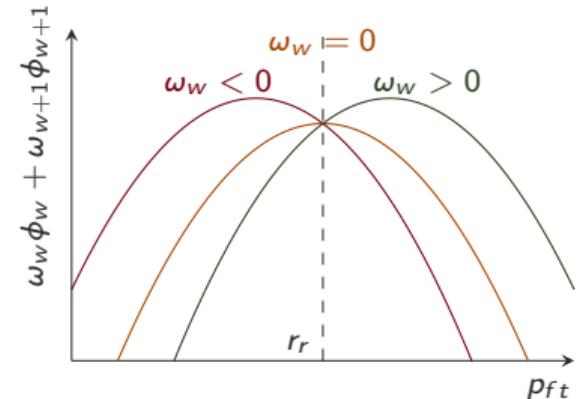
$$\phi(\mathbf{P}_{1:t-1}, p_{ft}) = \begin{bmatrix} \vdots \\ 1[p_{ft} = p_a] \\ \vdots \\ p_{ft} - r_r(\mathbf{P}_{1:t-1}) \\ (p_{ft} - r_r(\mathbf{P}_{1:t-1}))^2 \\ \vdots \end{bmatrix} \approx \begin{bmatrix} \vdots \\ \text{Context invariant price indicators} \\ \vdots \\ \text{Quadratic relative pricing rules} \\ \vdots \end{bmatrix}$$

- Price history reference points  $r_1, \dots, r_R$ :

- Most recent prices
  - Own
  - Competitor mean
  - Competitor minimum
- Recency weighted average prices
  - Own
  - Competitor mean
  - Competitor minimum

$$\hat{Q} = \omega' \cdot \phi = \begin{bmatrix} \vdots \\ \omega_w \\ \omega_{w+1} \\ \vdots \end{bmatrix}' \cdot \begin{bmatrix} p_{ft} - r_r(\mathbf{P}_{1:t-1}) \\ (p_{ft} - r_r(\mathbf{P}_{1:t-1}))^2 \\ \vdots \end{bmatrix}$$

$$(\omega_{w+1} < 0)$$



# Setup

Algorithm Hyperparameters

## Discounting

- $\delta$  : Discount rate
  - Governs intertemporal tradeoffs

$$Q(p_{ft} | \mathcal{P}_{1:t-1}) = \pi_{ft} + \delta \max_{p_{ft+1}} Q(p_{ft+1} | \mathcal{P}_{1:t})$$

## Exploration

- $\epsilon$  : Exploration rate
  - Determines how often the algorithm chooses a random price

$$p_{ft}(\mathcal{P}_{1:t-1}) = \begin{cases} \arg \max_{p_{ft}} Q(p_{ft} | \mathcal{P}_{1:t-1}) & \text{if: } \exp(-\frac{t}{\epsilon}) < U[0, 1] \\ p \in_R \mathcal{P} & \text{otherwise} \end{cases}$$

## Learning

- $\lambda$  : Learning rate
  - Determines how quickly old experience is replaced with new experience

$$\omega \leftarrow \omega + (\lambda \cdot e_b \cdot \phi)$$

## Price Memory

- $\rho$  : Recency weighting
  - Determines how much weight older prices receive in state representation

$$\bar{p}_t = \rho \overline{p_{t-1}} + (1 - \rho) p_t$$

# Setup

Baseline & Robustness

## Simulations

- Simulations:  $S = [1 \quad 1000]$
- Time Periods:  $T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$

## Strategic Environment

- Actions:  $A = [21 \quad 41]$
- Firms:  $F = [2 \quad 3 \quad 4]$
- Differentiation:  $\mu = [0.02 \quad 0.04 \quad 0.06 \quad 0.08 \quad 0.10]$

## Algorithm Parameters

- Discount Rate:  $\delta = [0.2 \quad 0.4 \quad 0.6 \quad 0.8 \quad 0.9]$
- Exploration Intensity:  $\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$
- Learning Rate:  $\lambda = [0.08 \quad 0.12 \quad 0.16 \quad 0.20 \quad 0.24]$
- Price Memory:  $\rho = [0.00 \quad 0.05 \quad 0.10 \quad 0.25 \quad 0.50]$

# Results

Baseline - Single Simulation

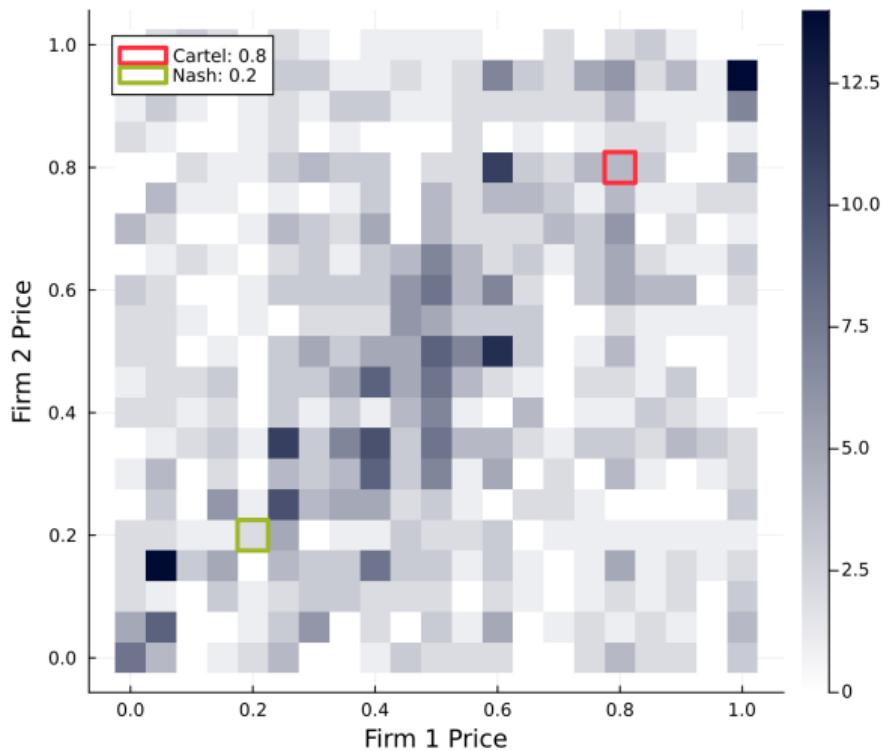
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Algorithms explore, then price match, then collude



# Results

Baseline - Single Simulation

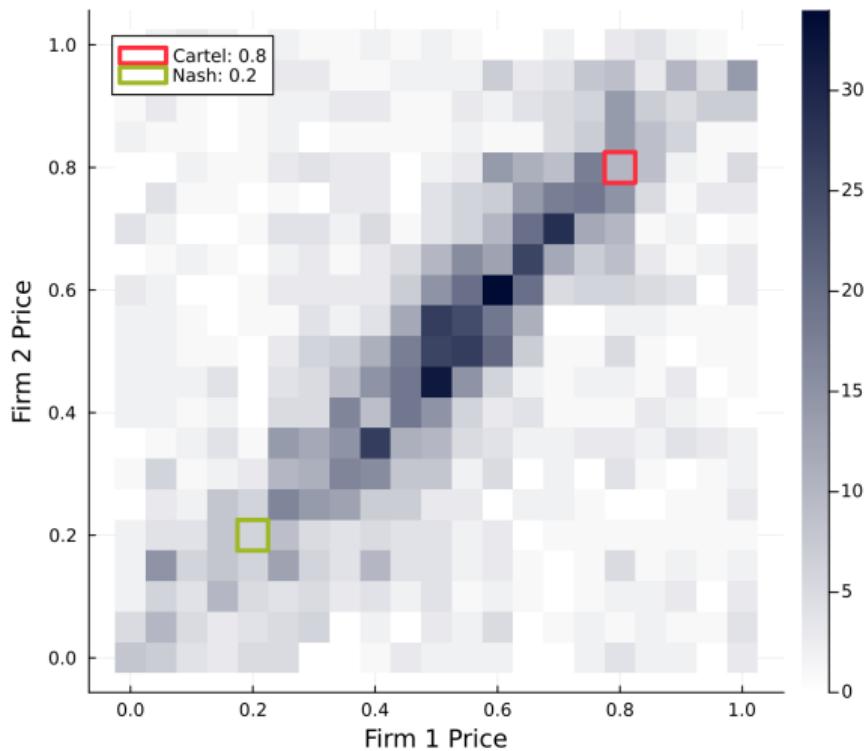
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Algorithms explore, then price match, then collude



# Results

Baseline - Single Simulation

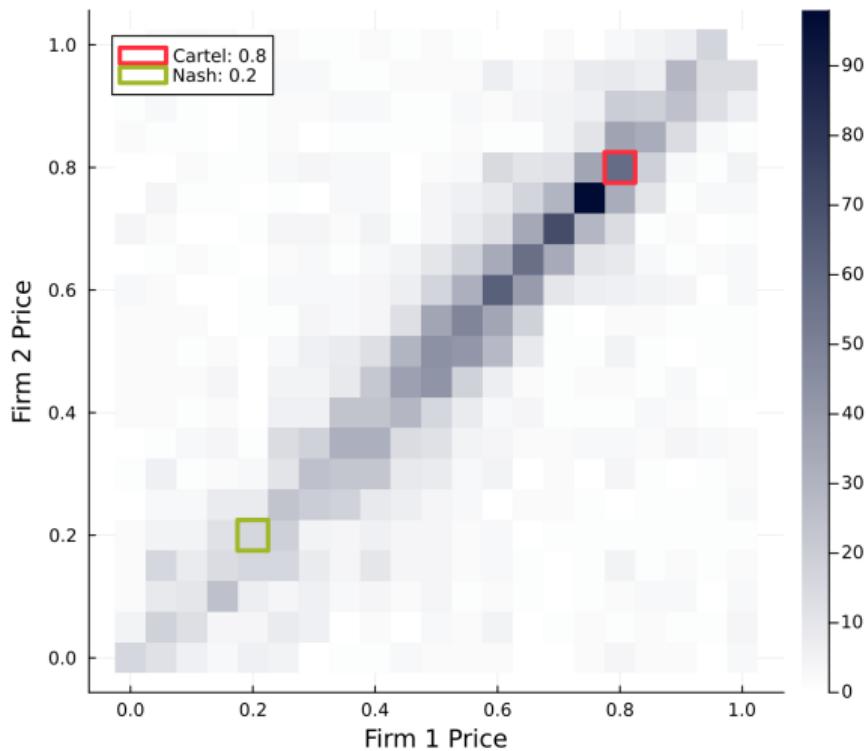
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Algorithms explore, then price match, then collude



# Results

Baseline - Single Simulation

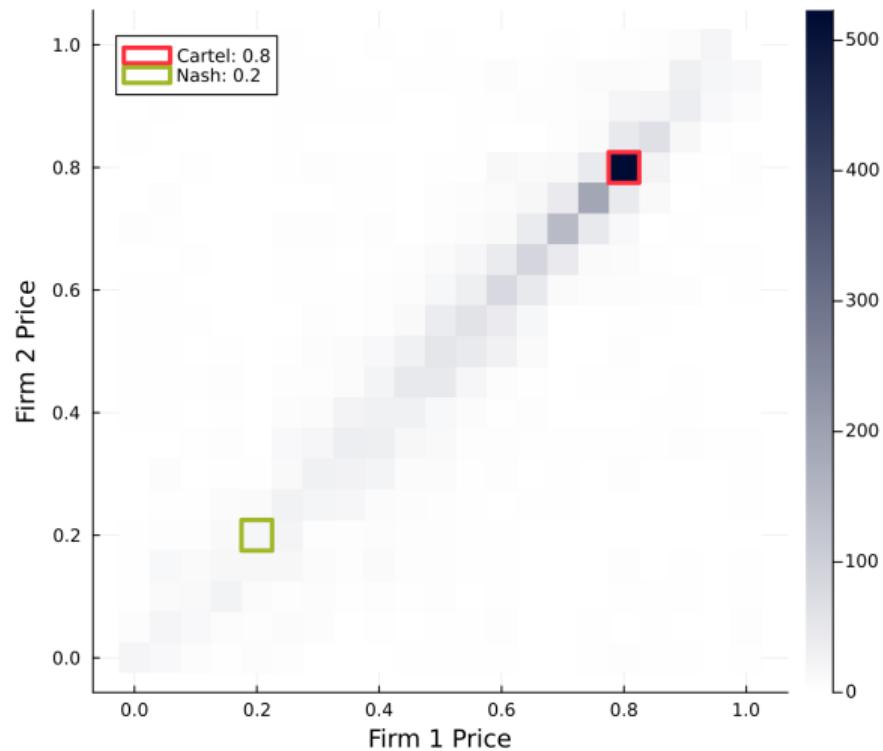
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Algorithms explore, then price match, then collude



# Results

Baseline - Single Simulation

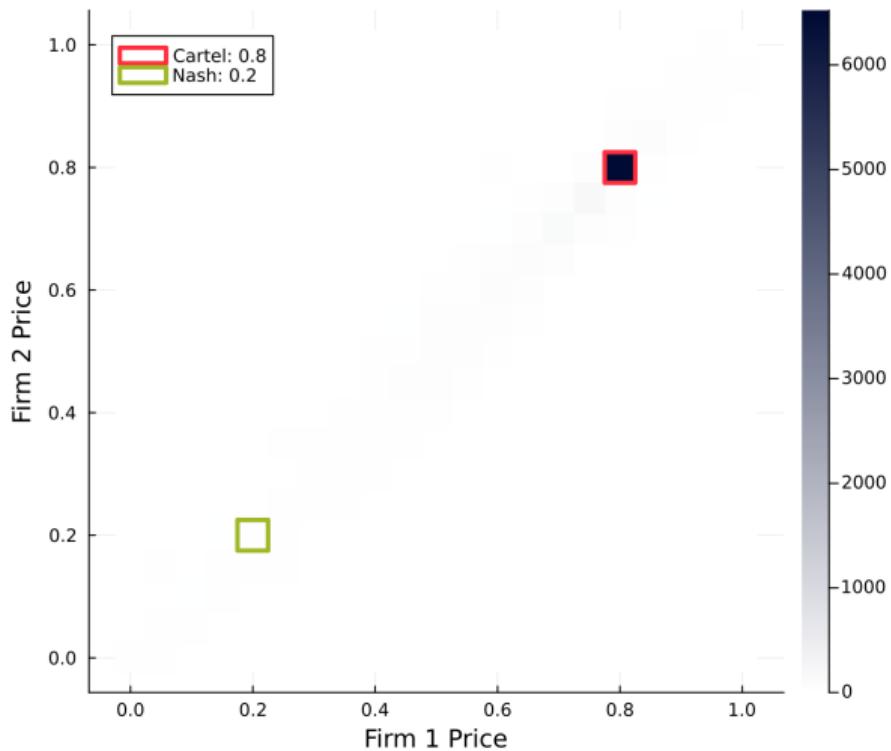
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Algorithms explore, then price match, then collude



# Results

Baseline - Empirical Distribution

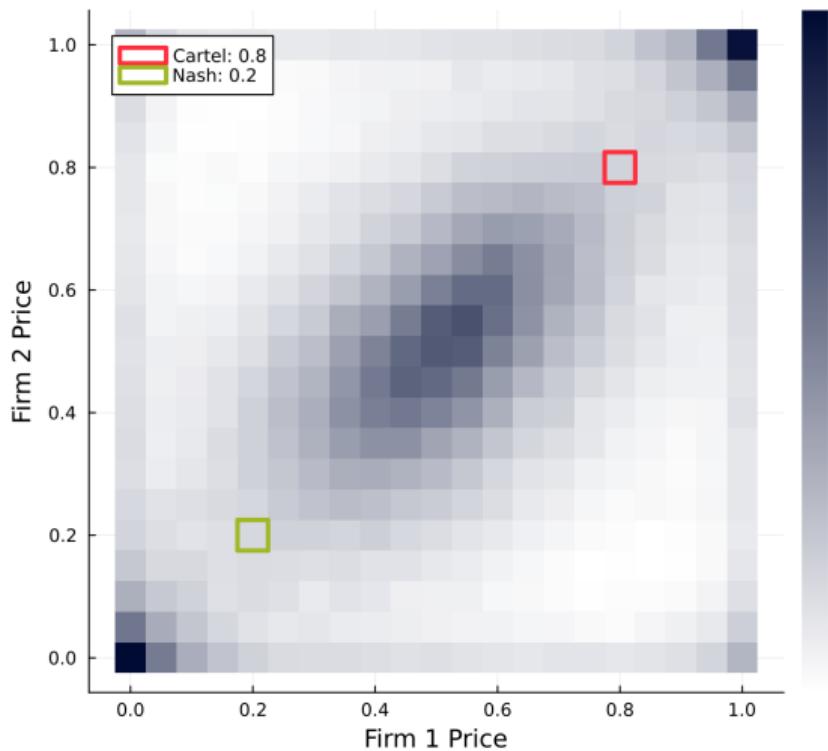
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Across 1000 simulations, similar action profiles
- Algorithms explore, then price match, then collude



# Results

Baseline - Empirical Distribution

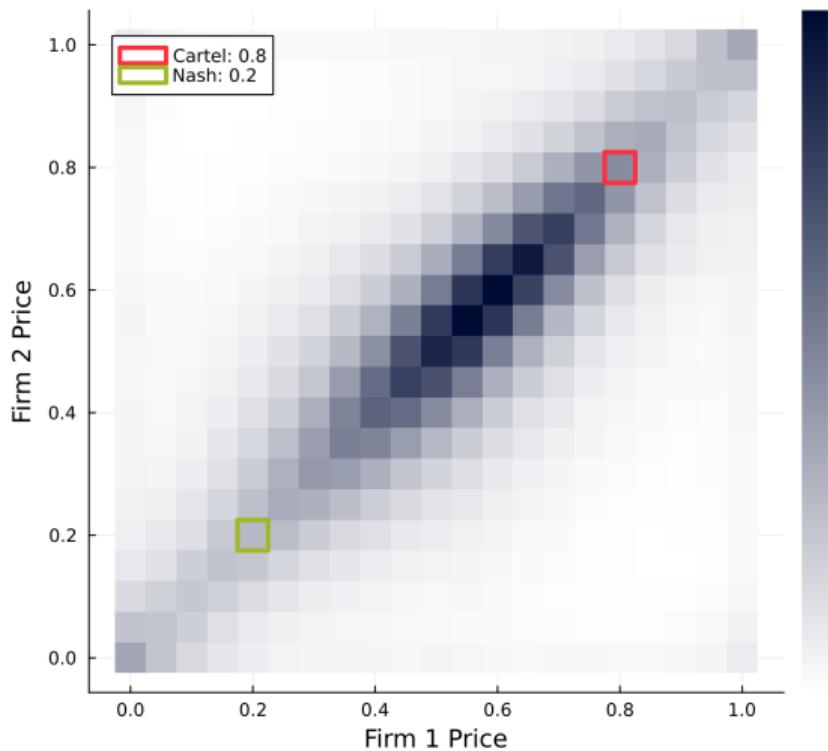
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Across 1000 simulations, similar action profiles
- Algorithms explore, then price match, then collude



# Results

Baseline - Empirical Distribution

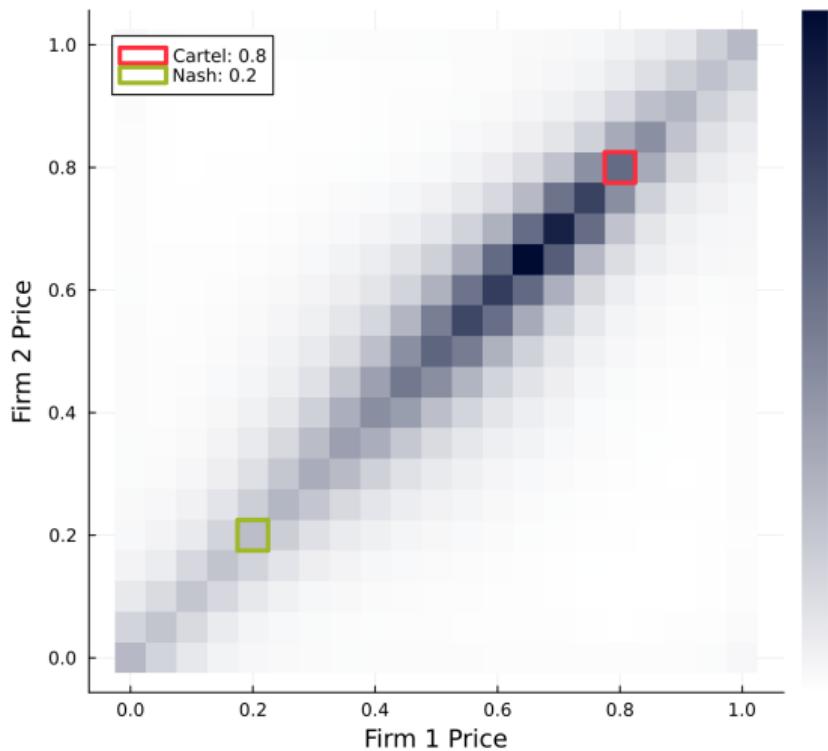
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Across 1000 simulations, similar action profiles
- Algorithms explore, then price match, then collude



# Results

Baseline - Empirical Distribution

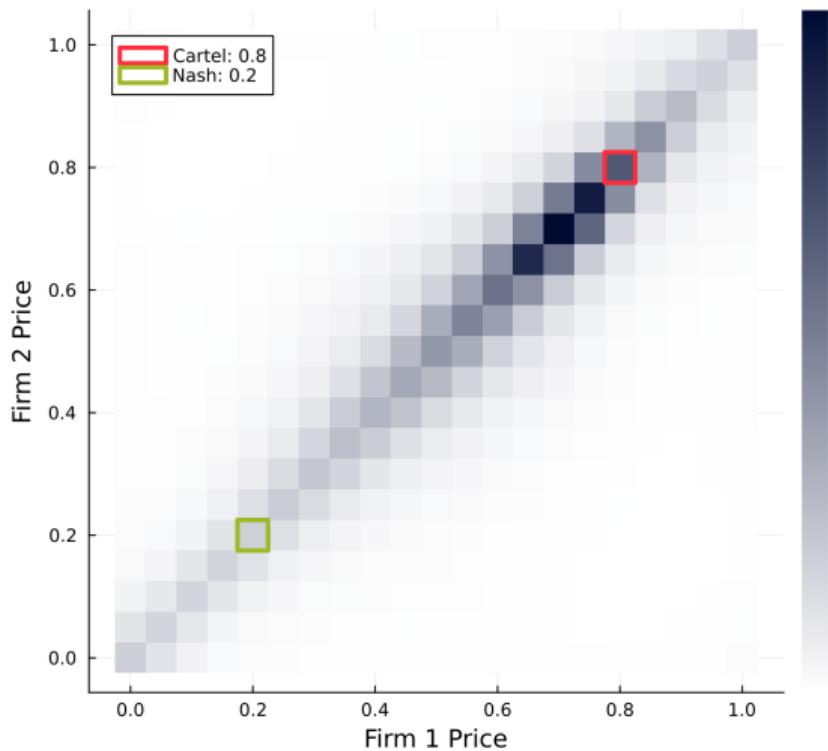
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Across 1000 simulations, similar action profiles
- Algorithms explore, then price match, then collude



# Results

Baseline - Empirical Distribution

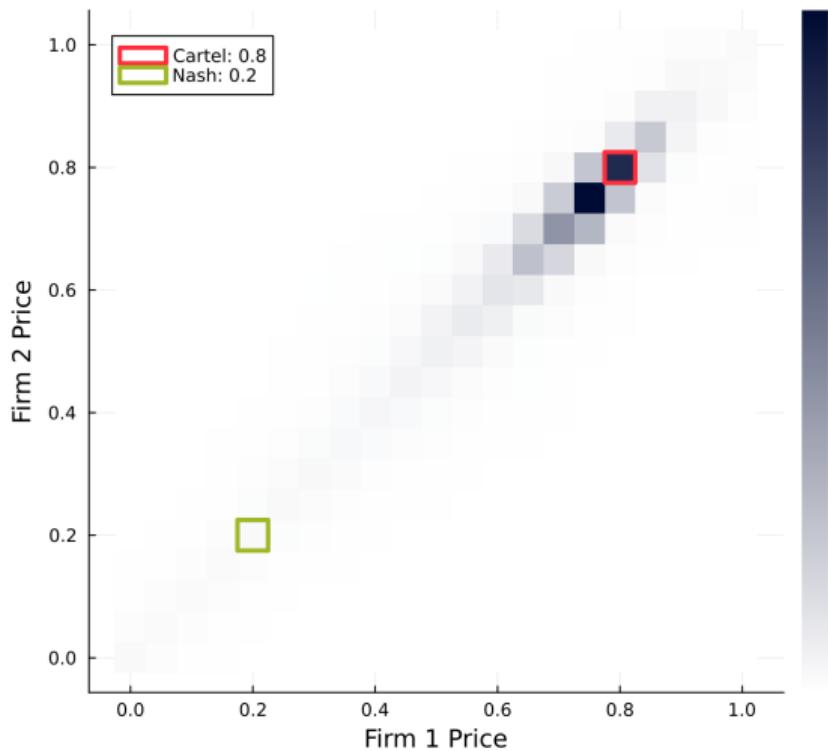
Simulations:

$$S = [1 \quad 1000]$$

Time Periods:

$$T = [1000 \quad 2000 \quad 3000 \quad 4000 \quad 10000]$$

- Across 1000 simulations, similar action profiles
- Algorithms explore, then price match, then collude



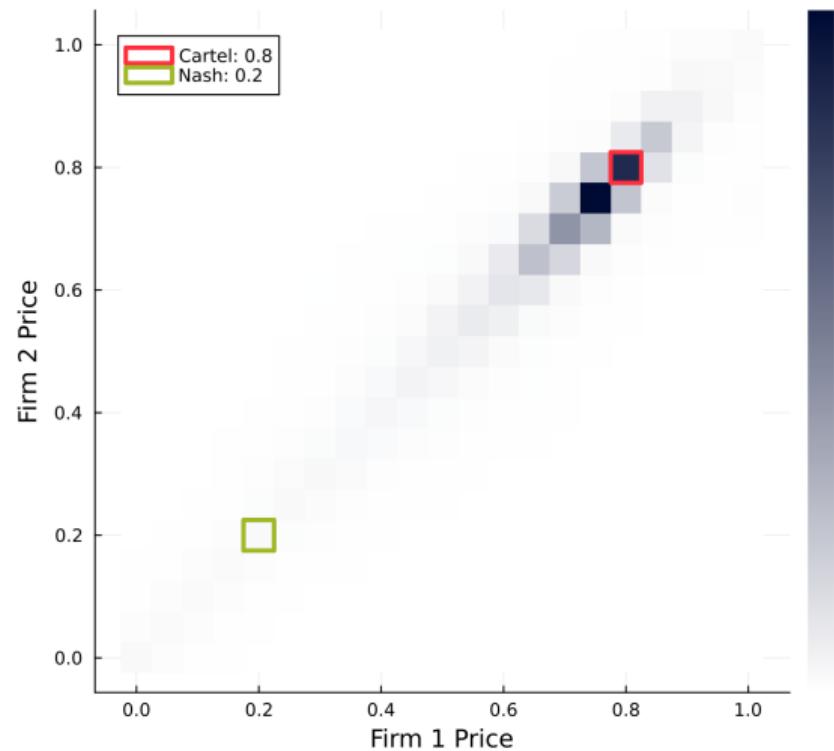
# Results

Robustness - Algorithm Parametrization - Discounting

Discount Factor:

$$\delta = [0.20 \quad 0.50 \quad 0.80 \quad 0.85 \quad 0.90]$$

- Sufficiently high discount factor required to sustain collusion
  - Consistent with theory on critical discount factors



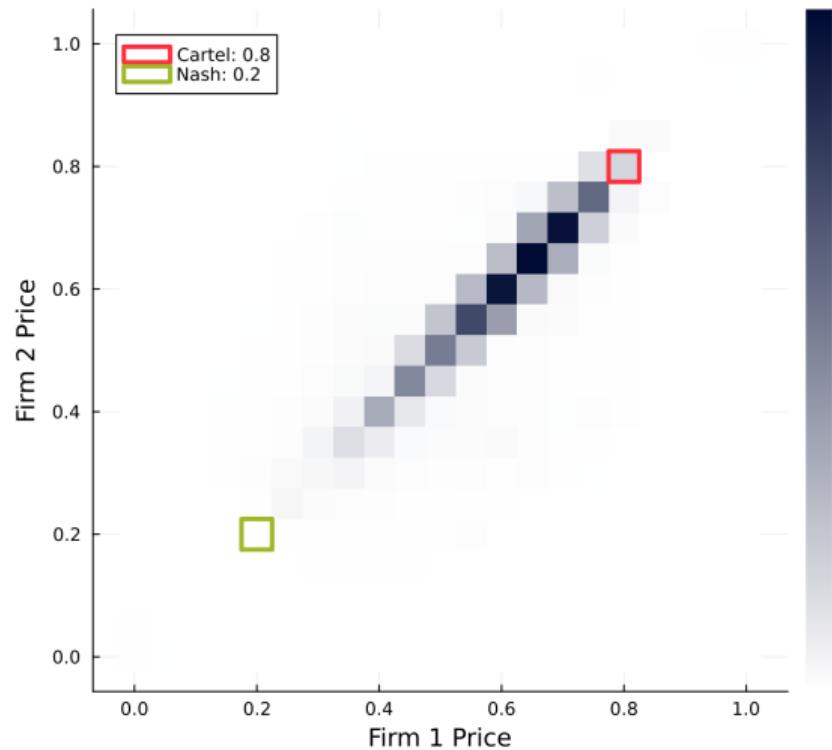
# Results

Robustness - Algorithm Parametrization - Discounting

Discount Factor:

$$\delta = [0.20 \quad 0.50 \quad 0.80 \quad 0.85 \quad 0.90]$$

- Sufficiently high discount factor required to sustain collusion
  - Consistent with theory on critical discount factors



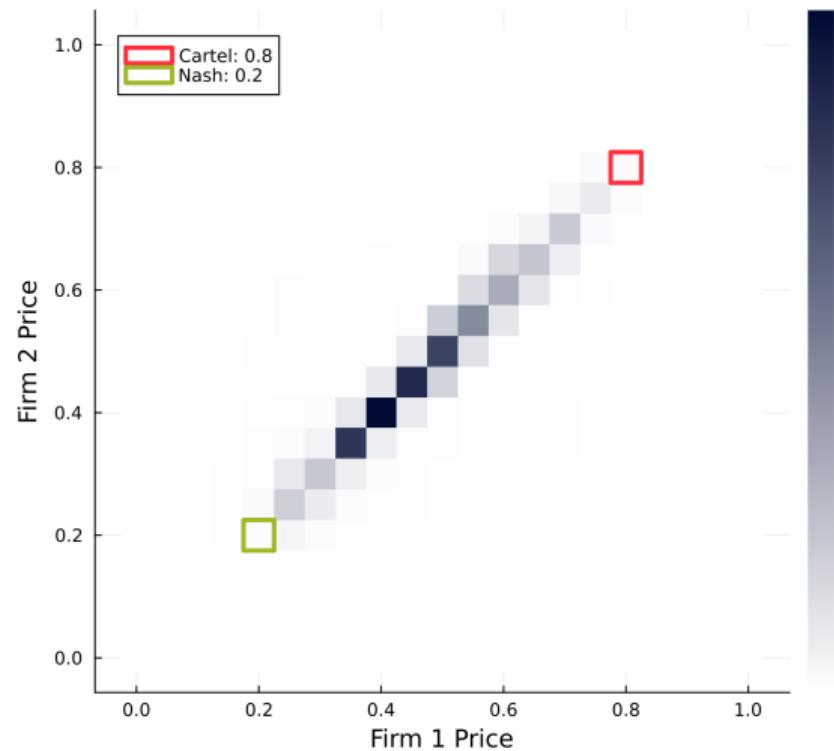
# Results

## Robustness - Algorithm Parametrization - Discounting

Discount Factor:

$$\delta = [0.20 \quad 0.50 \quad 0.80 \quad 0.85 \quad 0.90]$$

- Sufficiently high discount factor required to sustain collusion
  - Consistent with theory on critical discount factors



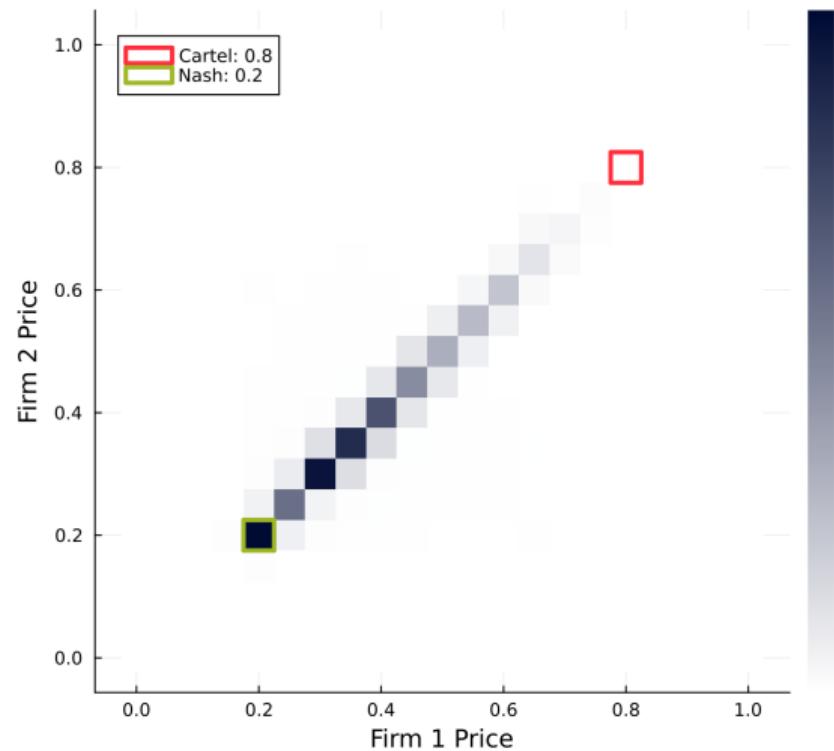
# Results

## Robustness - Algorithm Parametrization - Discounting

Discount Factor:

$$\delta = [0.20 \quad 0.50 \quad 0.80 \quad 0.85 \quad 0.90]$$

- Sufficiently high discount factor required to sustain collusion
  - Consistent with theory on critical discount factors



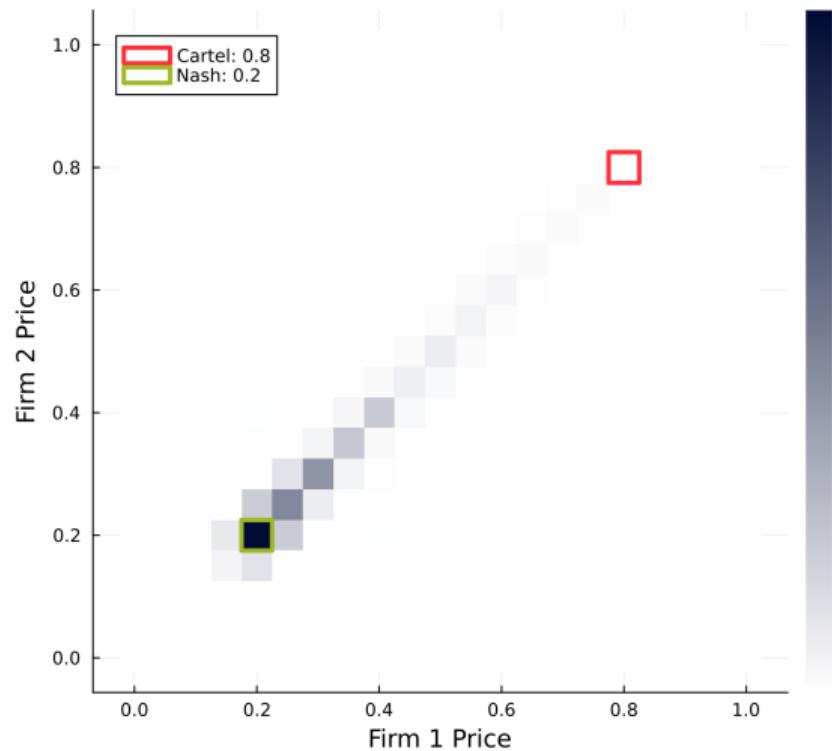
# Results

## Robustness - Algorithm Parametrization - Discounting

Discount Factor:

$$\delta = [0.20 \quad 0.50 \quad 0.80 \quad 0.85 \quad 0.90]$$

- Sufficiently high discount factor required to sustain collusion
  - Consistent with theory on critical discount factors



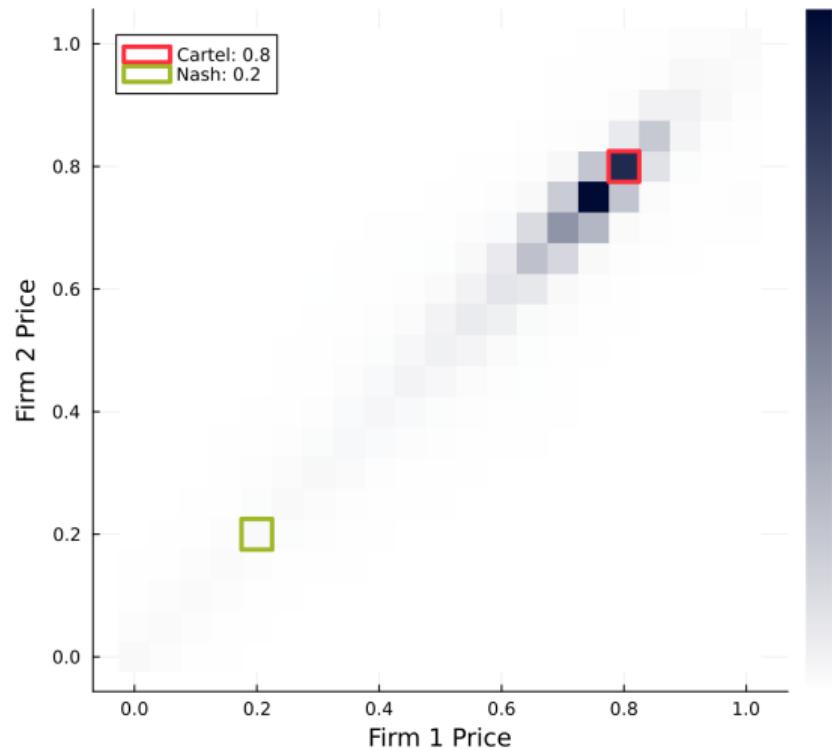
# Results

## Robustness - Algorithm Parametrization - Price Memory

Price Memory:

$$\rho = [0.00 \quad 0.05 \quad 0.10 \quad 0.25 \quad 0.50]$$

- By design, algorithms intermittently deviate from collusive prices
- Collusive stability requires distinguishing temporary and sustained deviation
  - Requires long-memory state representation



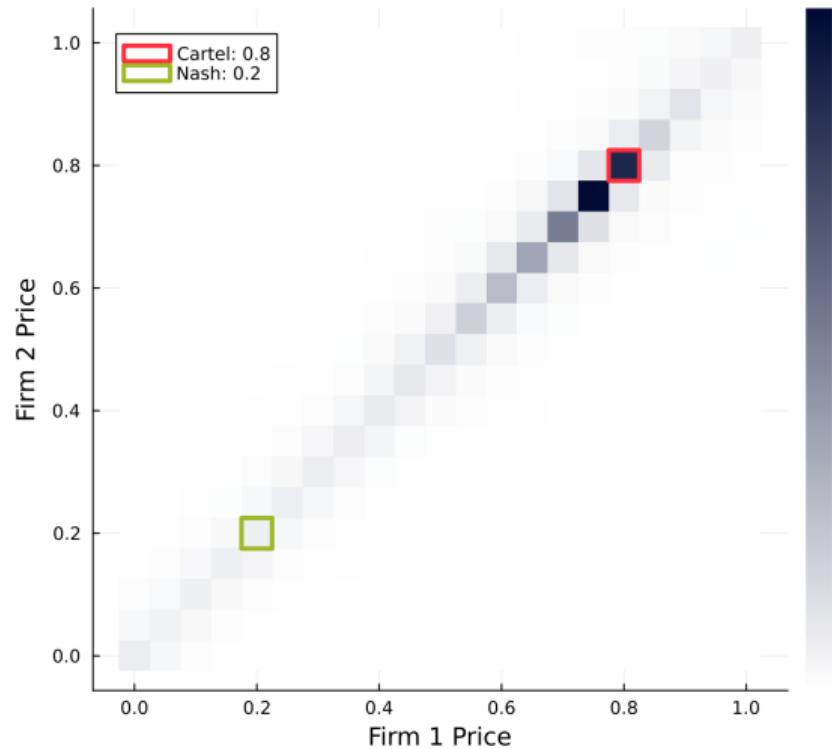
# Results

## Robustness - Algorithm Parametrization - Price Memory

Price Memory:

$$\rho = [0.00 \quad 0.05 \quad 0.10 \quad 0.25 \quad 0.50]$$

- By design, algorithms intermittently deviate from collusive prices
- Collusive stability requires distinguishing temporary and sustained deviation
  - Requires long-memory state representation



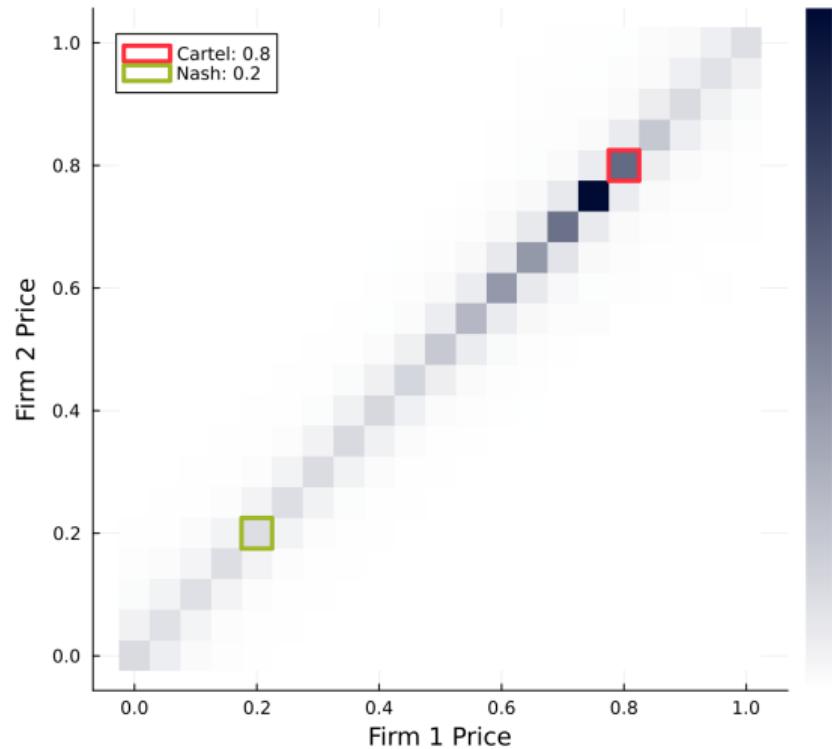
# Results

## Robustness - Algorithm Parametrization - Price Memory

Price Memory:

$$\rho = [0.00 \quad 0.05 \quad 0.10 \quad 0.25 \quad 0.50]$$

- By design, algorithms intermittently deviate from collusive prices
- Collusive stability requires distinguishing temporary and sustained deviation
  - Requires long-memory state representation



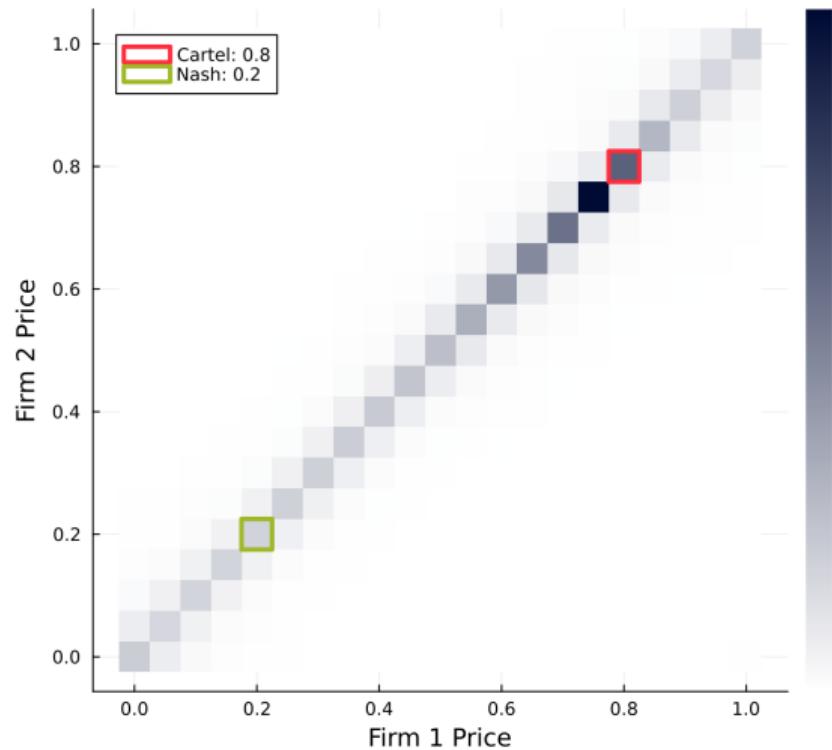
# Results

## Robustness - Algorithm Parametrization - Price Memory

Price Memory:

$$\rho = [0.00 \quad 0.05 \quad 0.10 \quad 0.25 \quad 0.50]$$

- By design, algorithms intermittently deviate from collusive prices
- Collusive stability requires distinguishing temporary and sustained deviation
  - Requires long-memory state representation



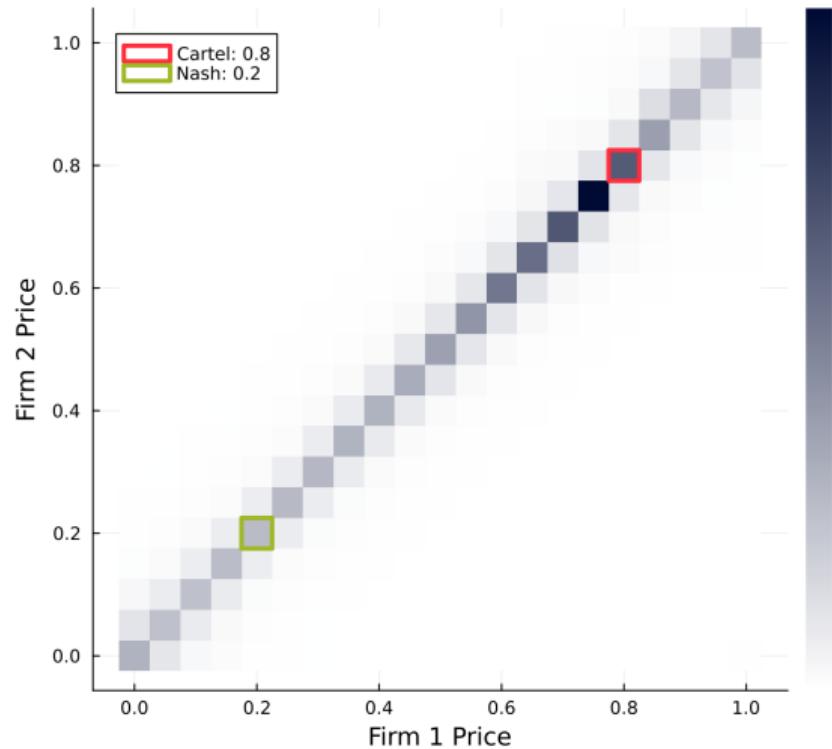
# Results

## Robustness - Algorithm Parametrization - Price Memory

Price Memory:

$$\rho = [0.00 \quad 0.05 \quad 0.10 \quad 0.25 \quad 0.50]$$

- By design, algorithms intermittently deviate from collusive prices
- Collusive stability requires distinguishing temporary and sustained deviation
  - Requires long-memory state representation



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

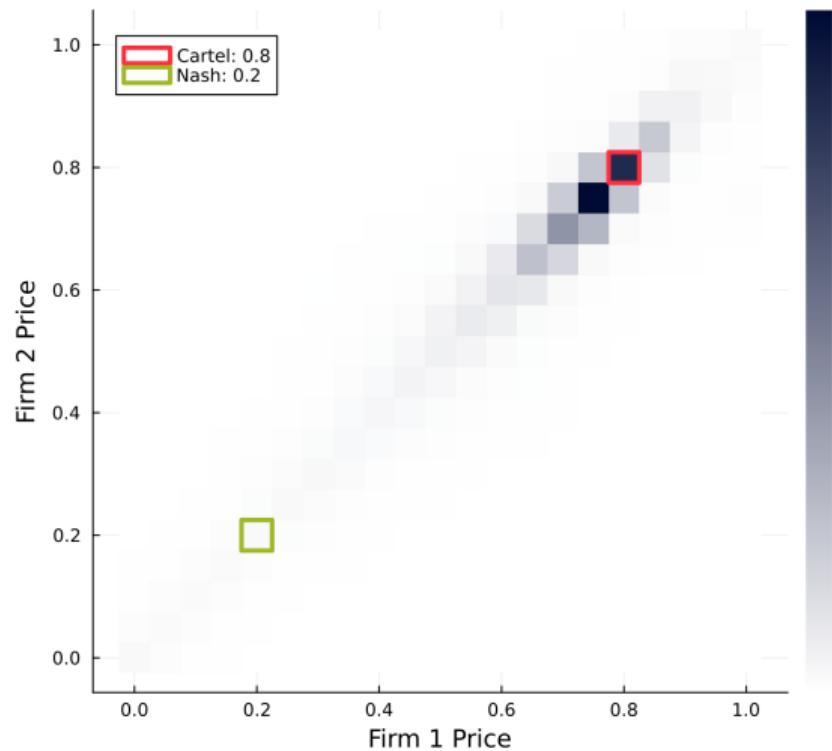
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.18 \quad 0.20]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

## Robustness - Algorithm Parametrization - Exploration & Learning

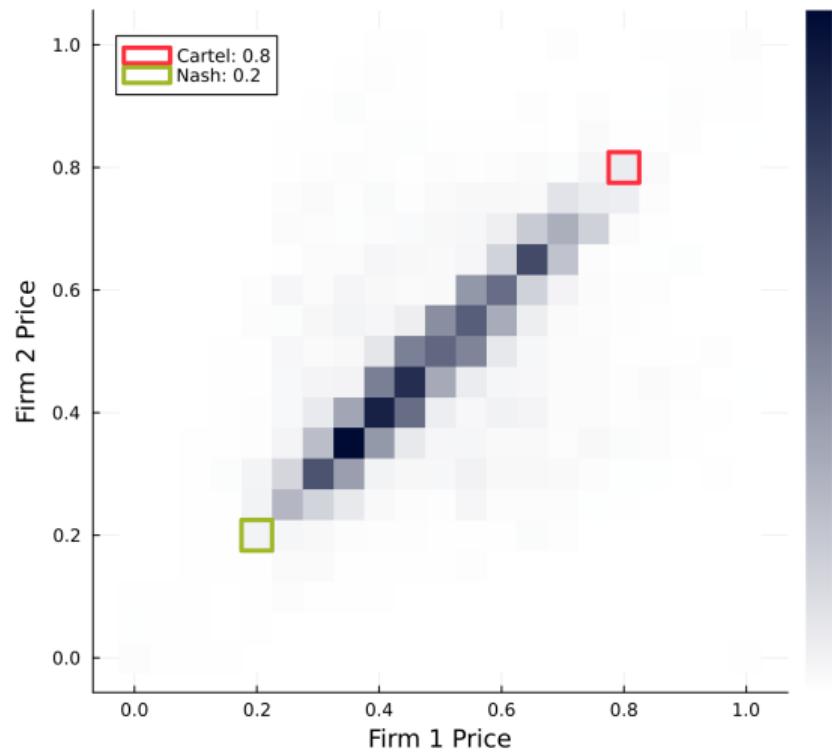
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.18 \quad 0.20]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

## Robustness - Algorithm Parametrization - Exploration & Learning

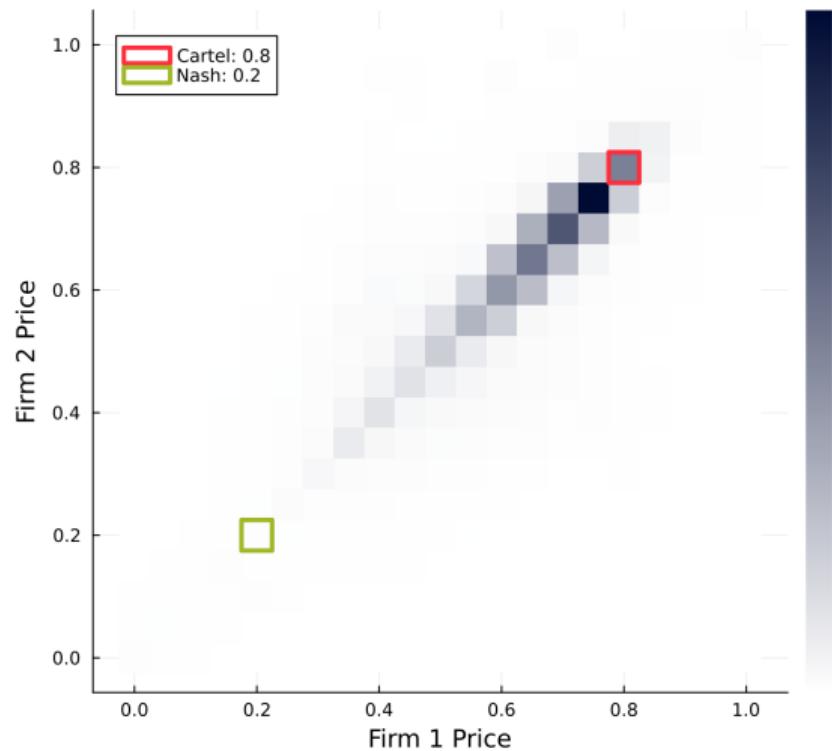
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.18 \quad 0.20]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

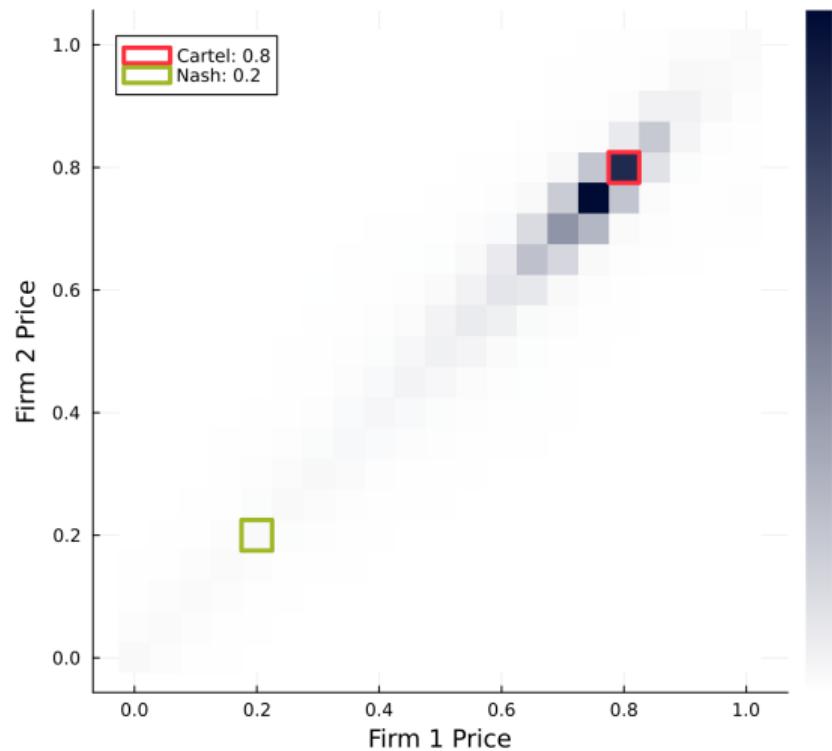
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.18 \quad 0.20]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

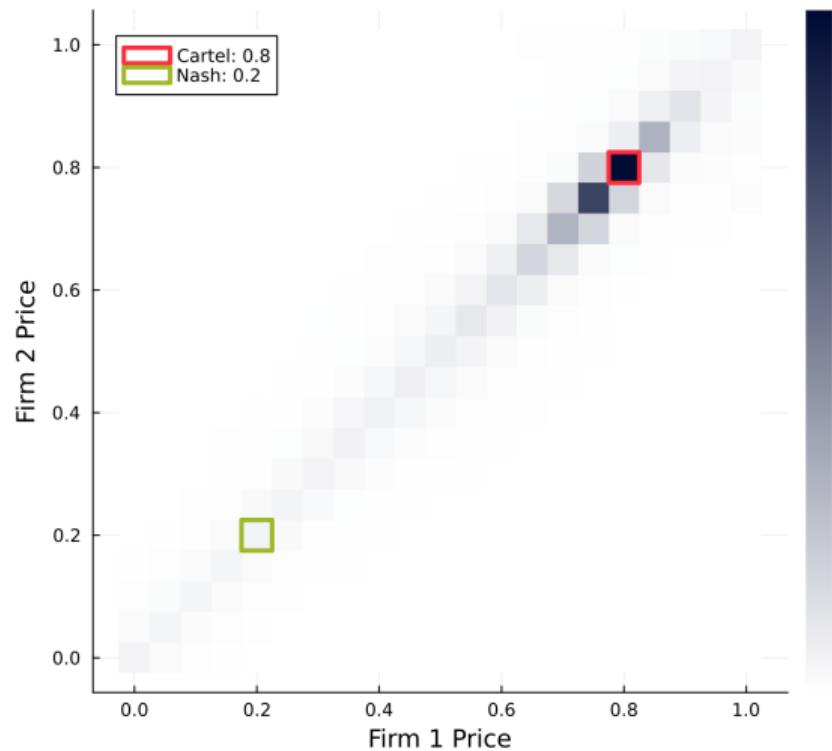
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.18 \quad 0.20]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

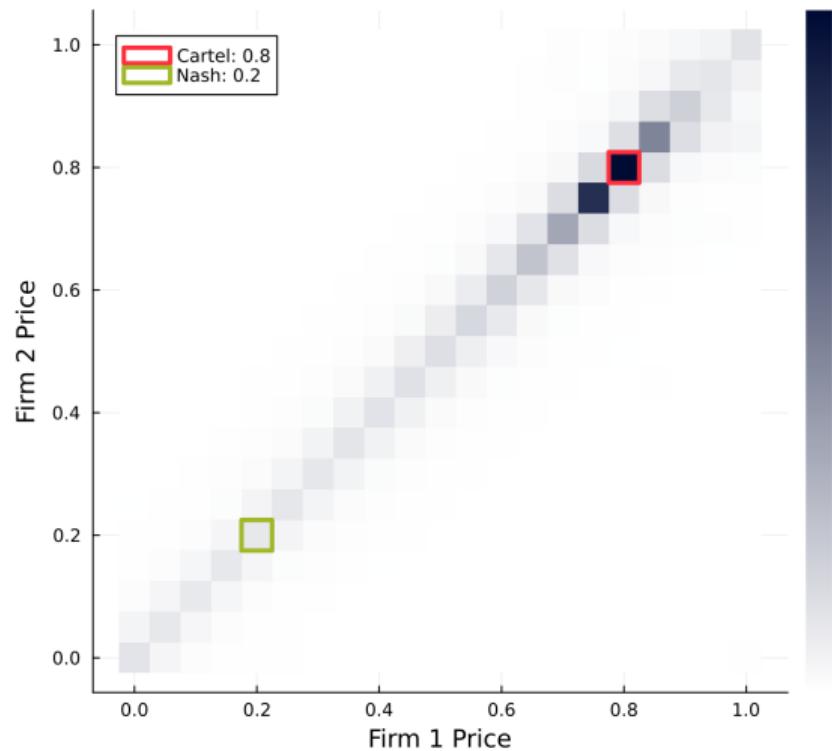
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.18 \quad 0.20]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

## Robustness - Algorithm Parametrization - Exploration & Learning

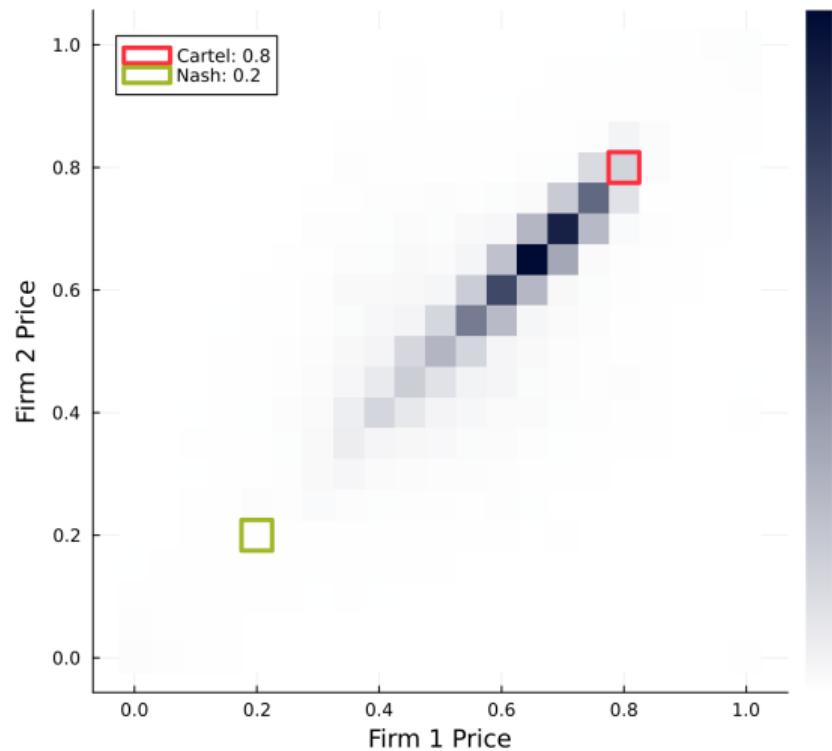
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

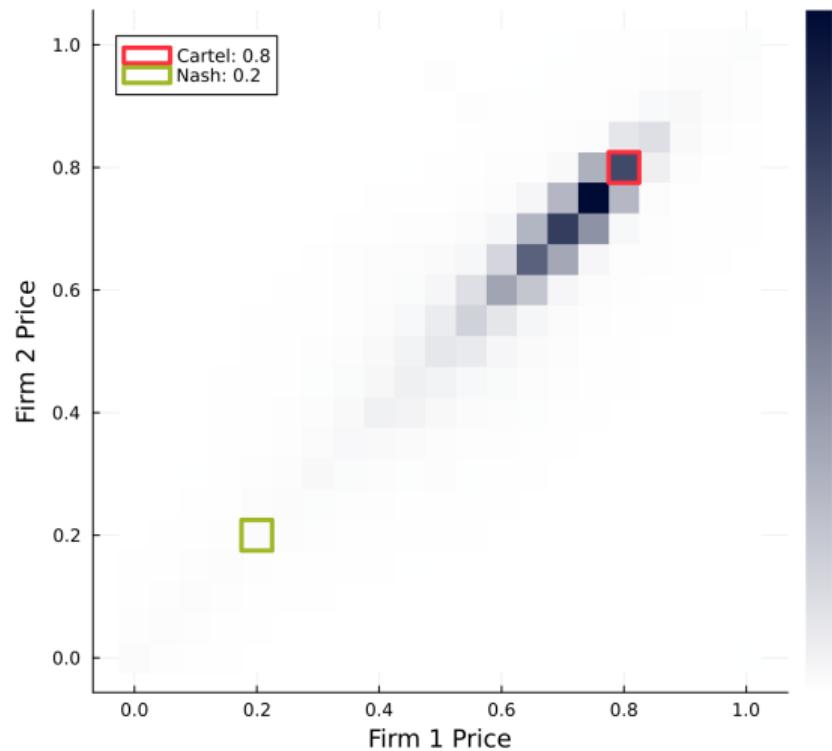
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

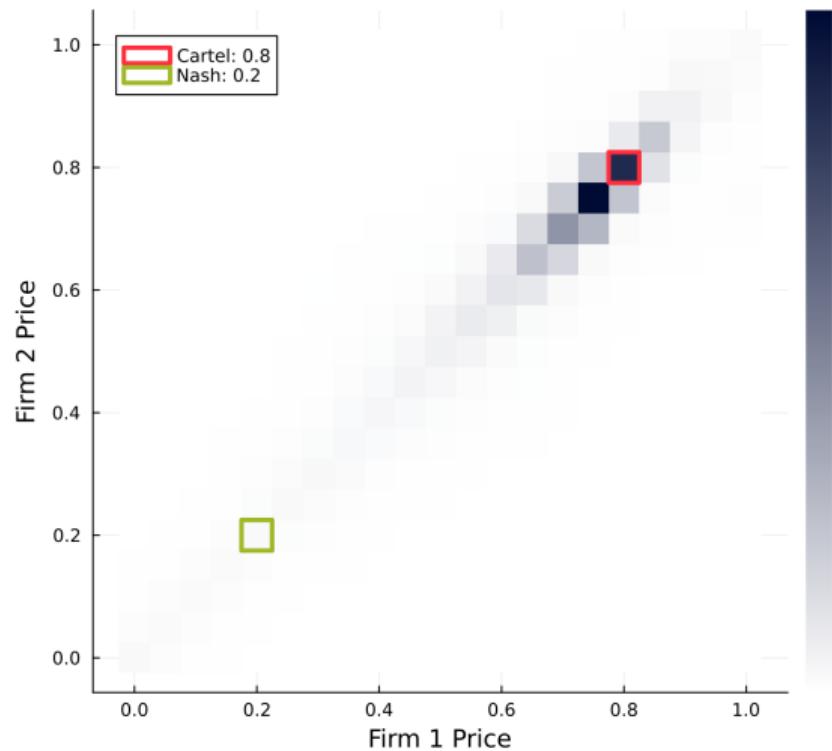
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

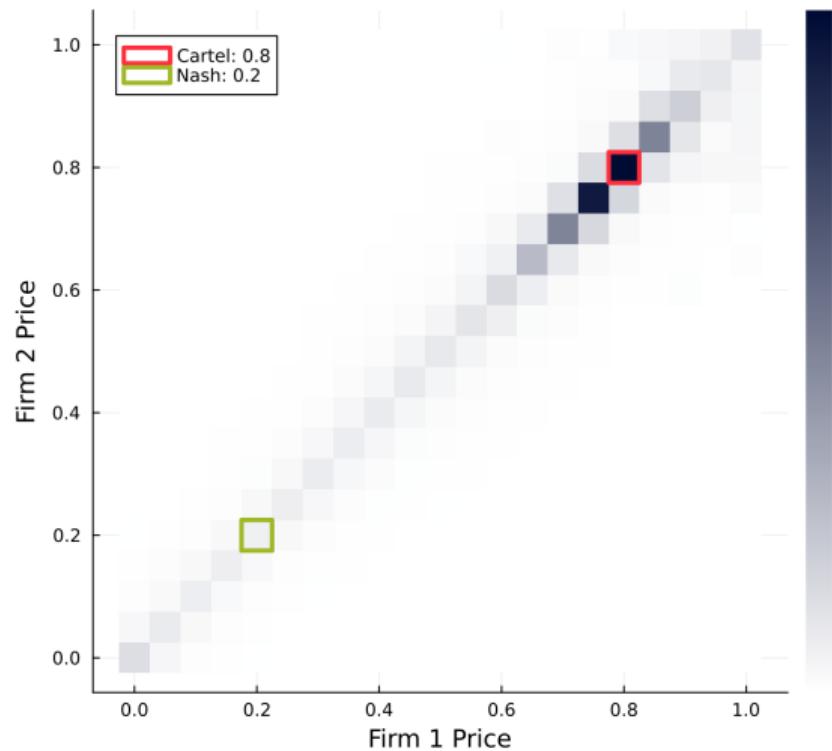
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

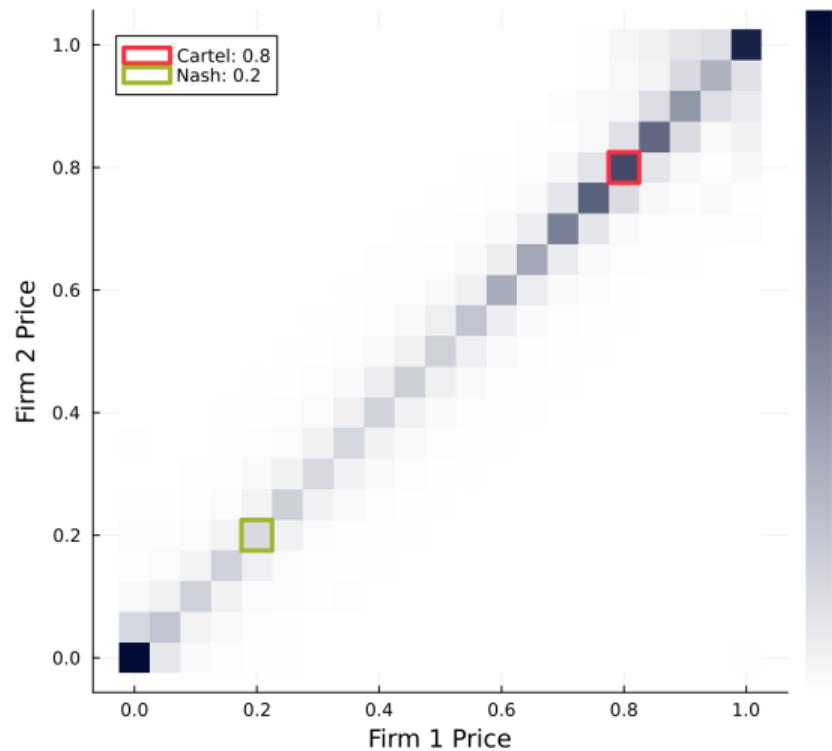
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

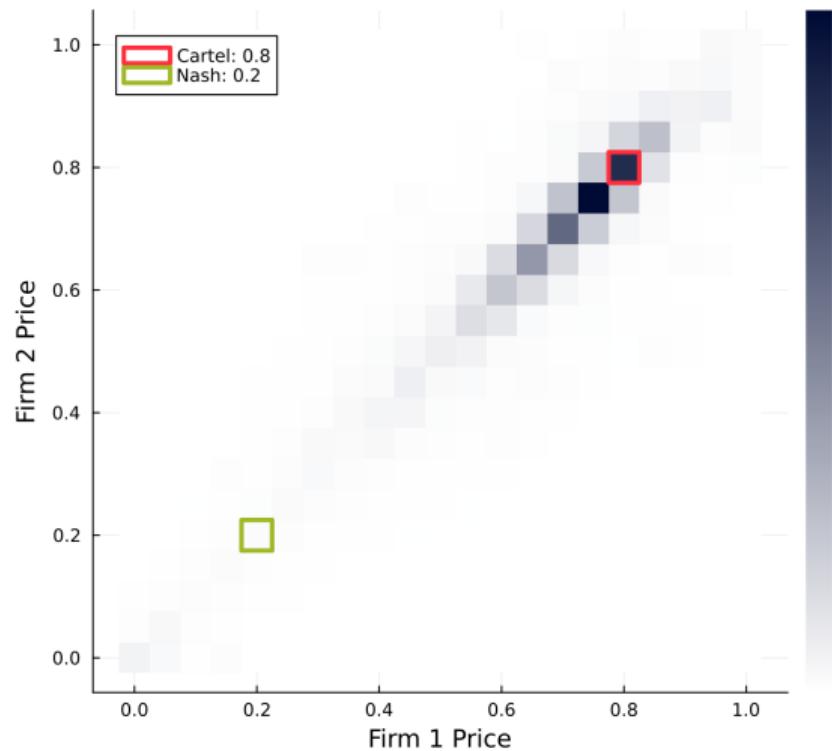
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

## Robustness - Algorithm Parametrization - Exploration & Learning

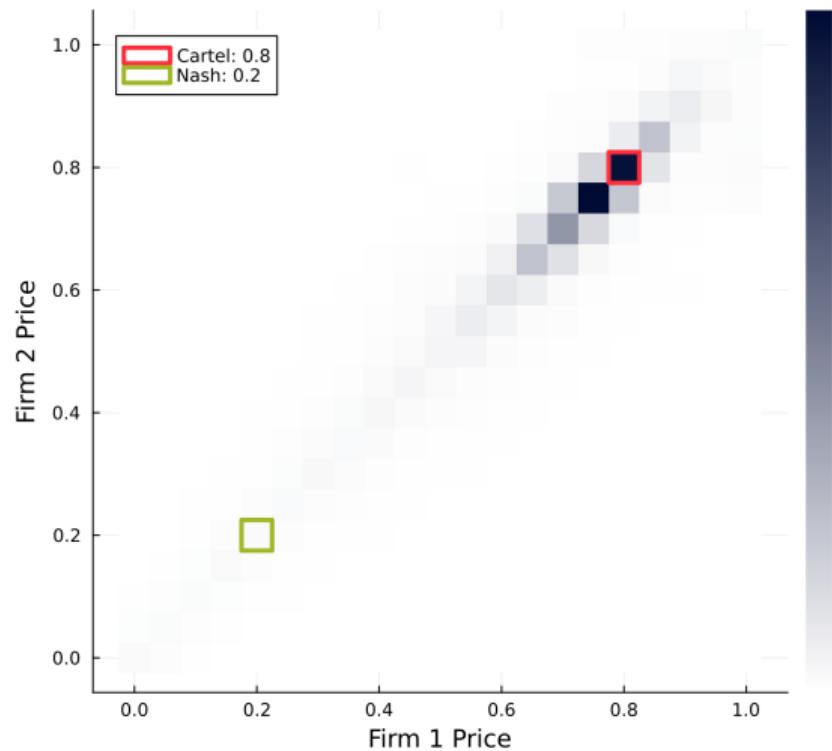
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

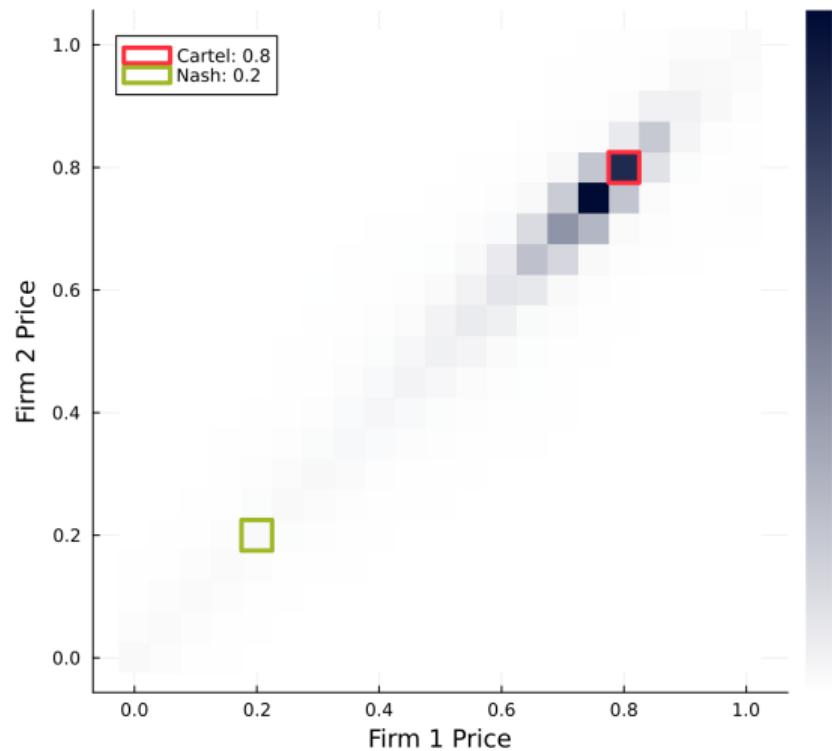
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

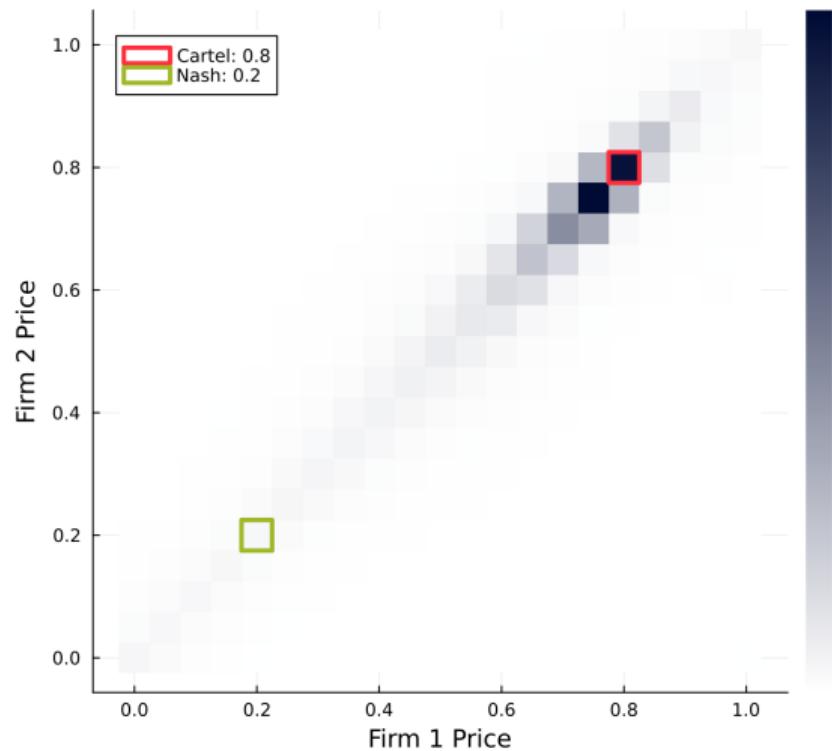
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning

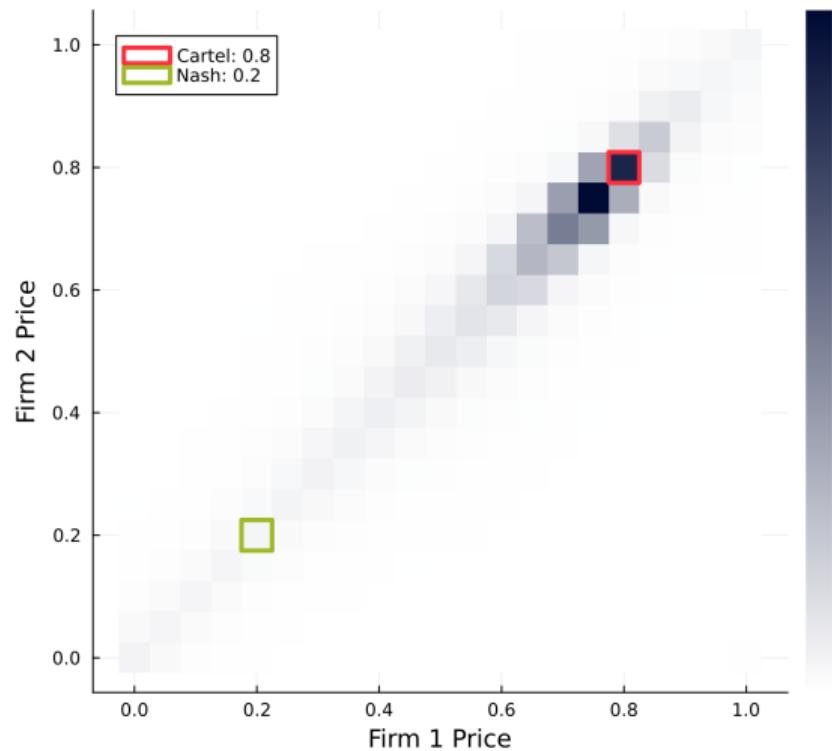
Exploration:

$$\varepsilon = [200 \quad 300 \quad 400 \quad 500 \quad 600]$$

Learning Rate:

$$\lambda = [0.12 \quad 0.14 \quad 0.16 \quad 0.20 \quad 0.24]$$

- Learning efficiency determinants:
  - Exploration frequency:  $\varepsilon$
  - Learning rate:  $\lambda$
- Efficient learning required to sustain collusion:
  - High  $\lambda$ , low  $\varepsilon$
  - Low  $\lambda$ , high  $\varepsilon$



# Results

Robustness - Algorithm Parametrization - Exploration & Learning - Asymmetric Firms

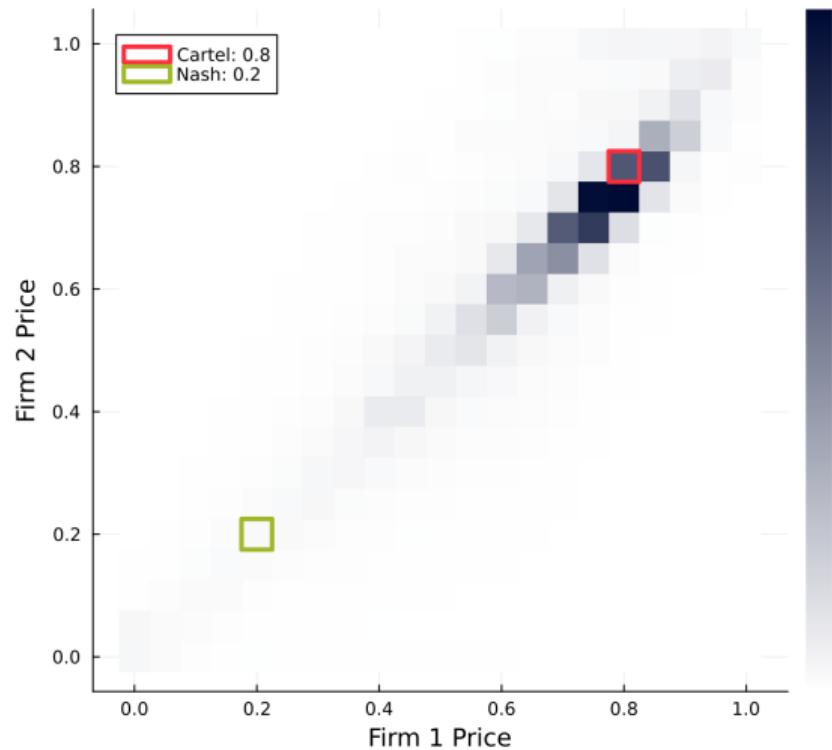
Exploration:

$$\begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \end{bmatrix} = \begin{bmatrix} 200 & 400 & 600 \\ 200 & 400 & 600 \end{bmatrix}$$

Learning Rate:

$$\begin{bmatrix} \lambda_1 \\ \lambda_2 \end{bmatrix} = \begin{bmatrix} 0.12 & 0.16 & 0.24 \\ 0.12 & 0.16 & 0.24 \end{bmatrix}$$

- Ex-ante strategic parametrization:
  - Faster learning algorithm exploits slower learning algorithm
- Asymmetric collusive outcomes still persist



# Results

Robustness - Algorithm Parametrization - Exploration & Learning - Asymmetric Firms

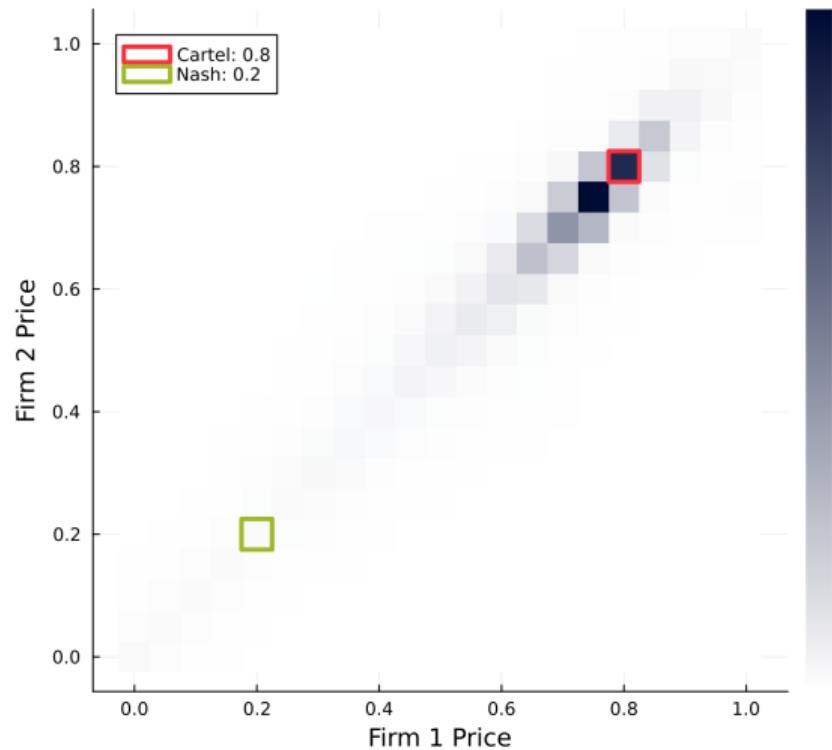
Exploration:

$$\begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \end{bmatrix} = \begin{bmatrix} 200 & 400 & 600 \\ 200 & 400 & 600 \end{bmatrix}$$

Learning Rate:

$$\begin{bmatrix} \lambda_1 \\ \lambda_2 \end{bmatrix} = \begin{bmatrix} 0.12 & 0.16 & 0.24 \\ 0.12 & 0.16 & 0.24 \end{bmatrix}$$

- Ex-ante strategic parametrization:
  - Faster learning algorithm exploits slower learning algorithm
- Asymmetric collusive outcomes still persist



# Results

Robustness - Algorithm Parametrization - Exploration & Learning - Asymmetric Firms

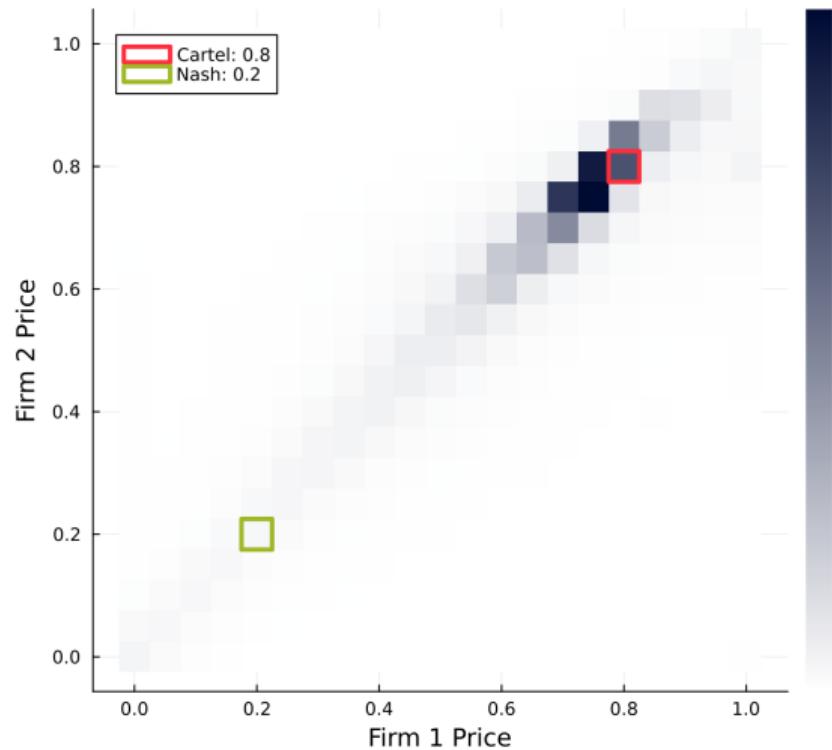
Exploration:

$$\begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \end{bmatrix} = \begin{bmatrix} 200 & 400 & 600 \\ 200 & 400 & 600 \end{bmatrix}$$

Learning Rate:

$$\begin{bmatrix} \lambda_1 \\ \lambda_2 \end{bmatrix} = \begin{bmatrix} 0.12 & 0.16 & 0.24 \\ 0.12 & 0.16 & 0.24 \end{bmatrix}$$

- Ex-ante strategic parametrization:
  - Faster learning algorithm exploits slower learning algorithm
- Asymmetric collusive outcomes still persist



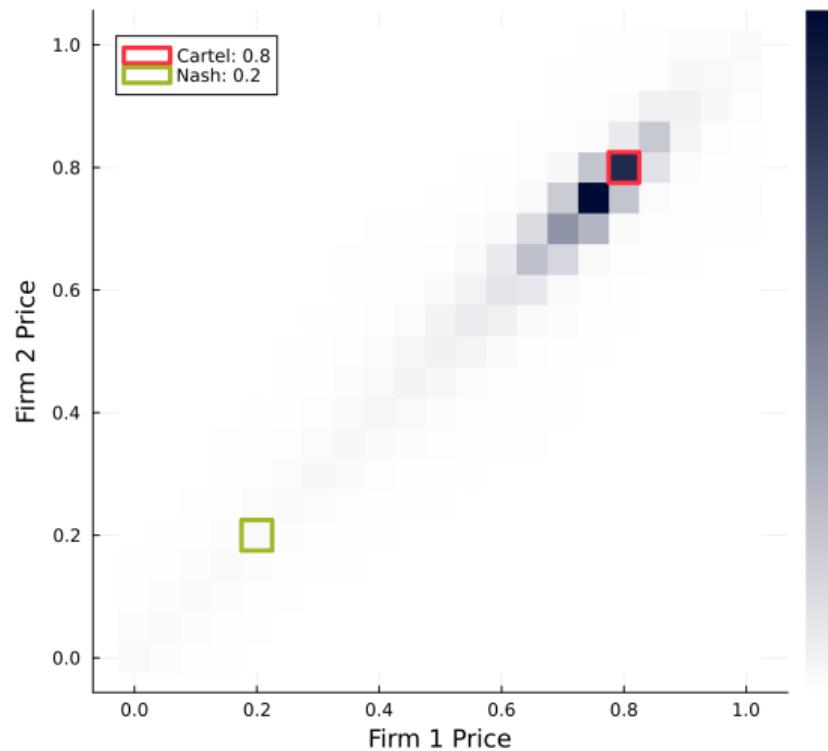
# Results

Robustness - Strategic Environment - Action Space

Action Space:

$$A = [21 \quad 41]$$

- Collusive outcomes robust to granularity of action space



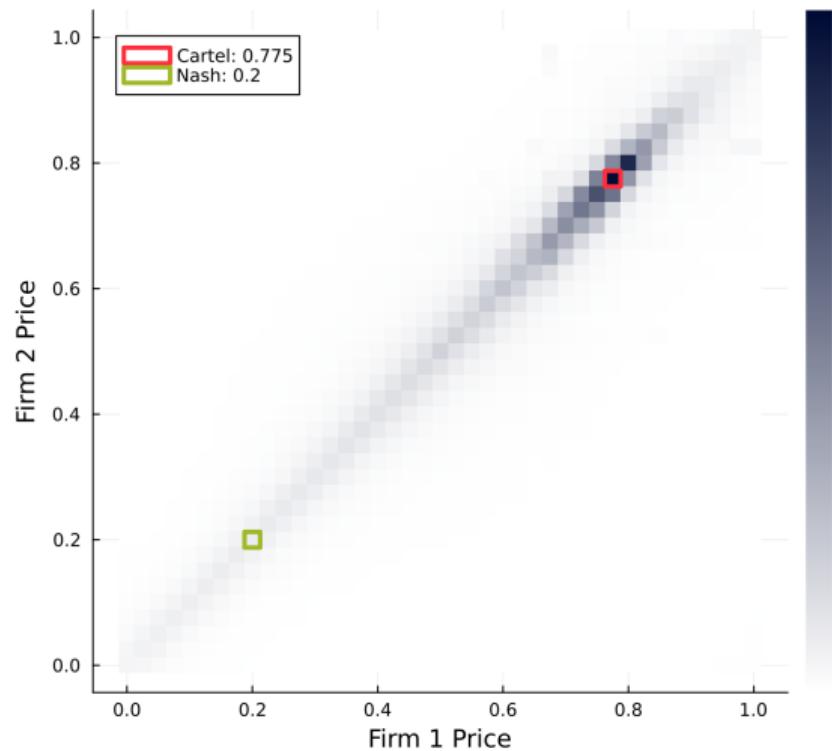
# Results

Robustness - Strategic Environment - Action Space

Action Space:

$$A = [21 \quad 41]$$

- Collusive outcomes robust to granularity of action space



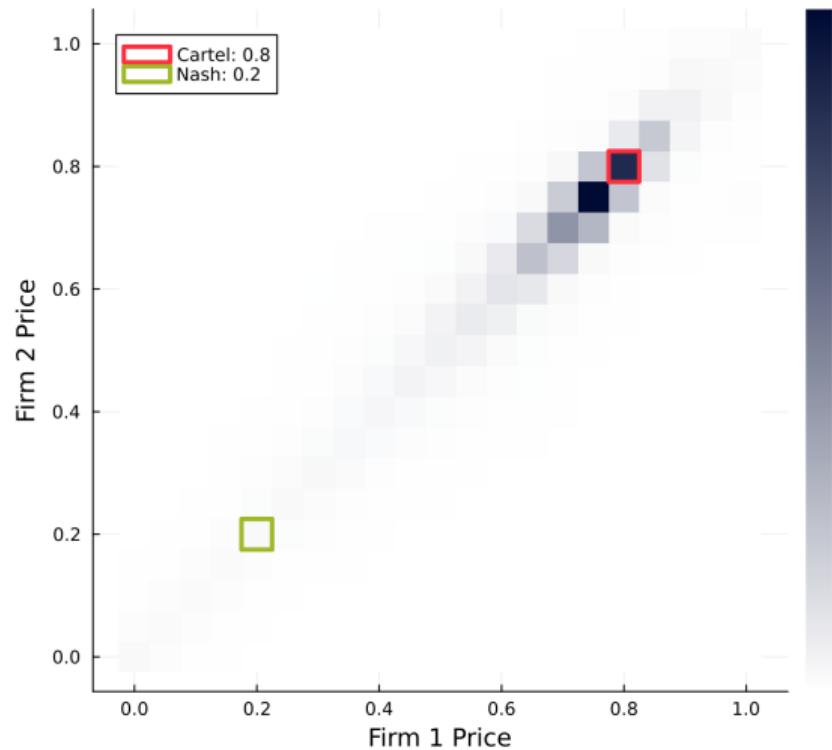
# Results

Robustness - Strategic Environment - Product Differentiation

Product Differentiation:

$$\mu = [0.02 \quad 0.04 \quad 0.06 \quad 0.08 \quad 0.10]$$

- Collusive outcomes robust to degree of product differentiation
- Product homogeneity increases price symmetry
  - Profits more sensitive to being undercut



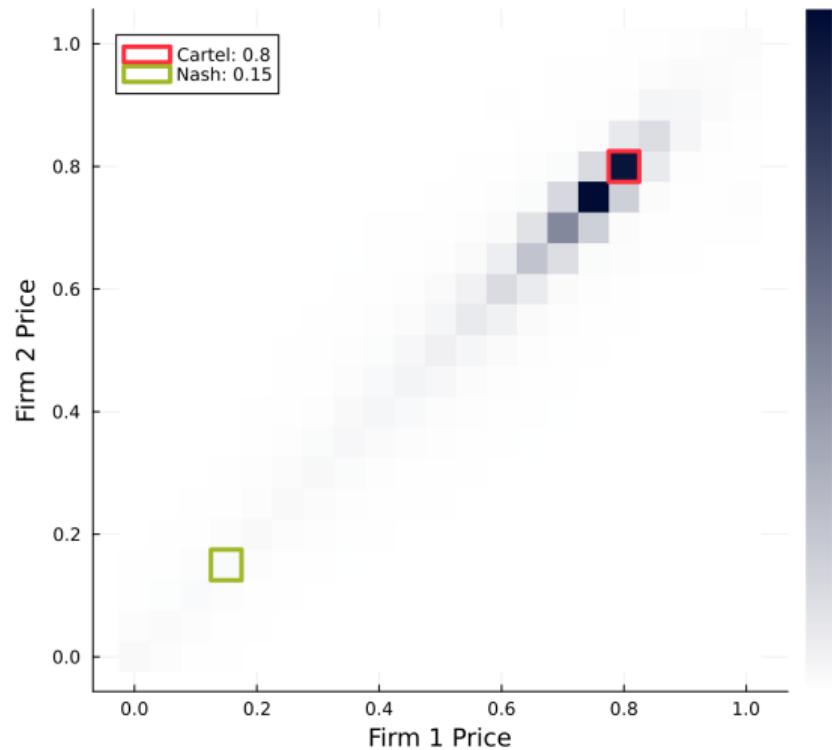
# Results

Robustness - Strategic Environment - Product Differentiation

Product Differentiation:

$$\mu = [0.02 \quad 0.04 \quad 0.06 \quad 0.08 \quad 0.10]$$

- Collusive outcomes robust to degree of product differentiation
- Product homogeneity increases price symmetry
  - Profits more sensitive to being undercut



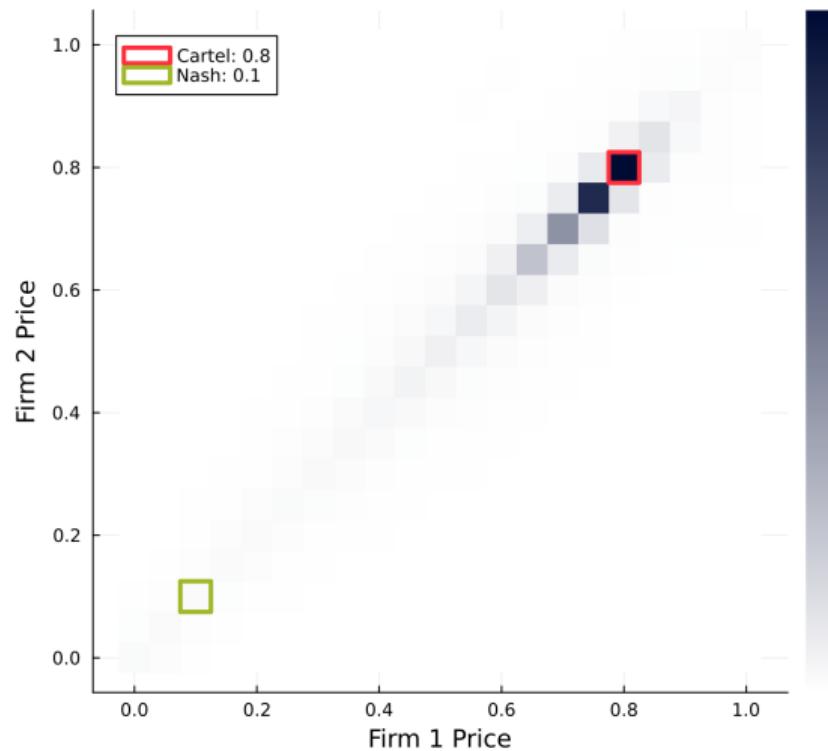
# Results

Robustness - Strategic Environment - Product Differentiation

Product Differentiation:

$$\mu = [0.02 \quad 0.04 \quad 0.06 \quad 0.08 \quad 0.10]$$

- Collusive outcomes robust to degree of product differentiation
- Product homogeneity increases price symmetry
  - Profits more sensitive to being undercut



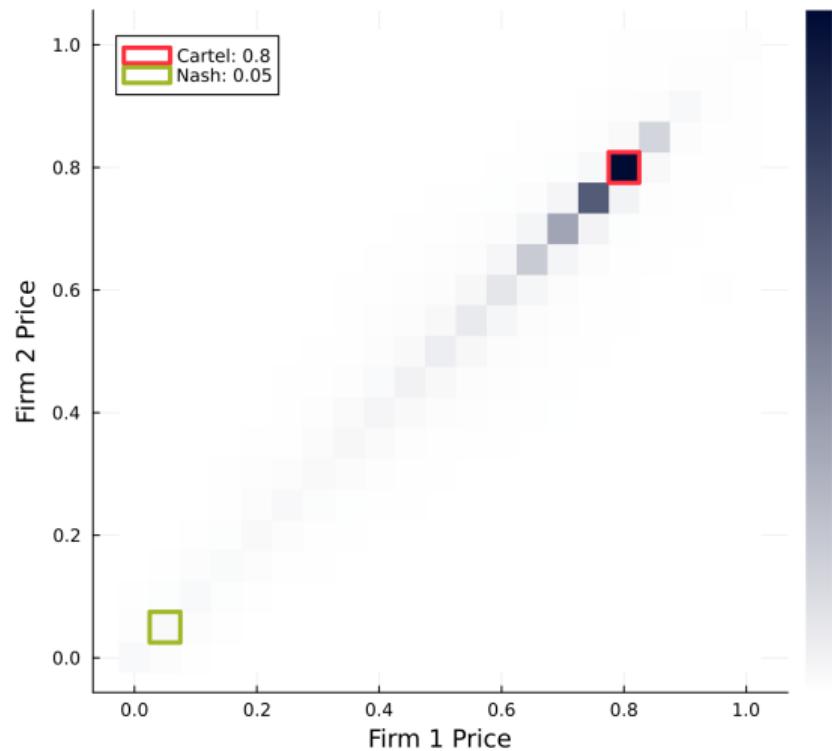
# Results

Robustness - Strategic Environment - Product Differentiation

Product Differentiation:

$$\mu = [0.02 \quad 0.04 \quad 0.06 \quad 0.08 \quad 0.10]$$

- Collusive outcomes robust to degree of product differentiation
- Product homogeneity increases price symmetry
  - Profits more sensitive to being undercut



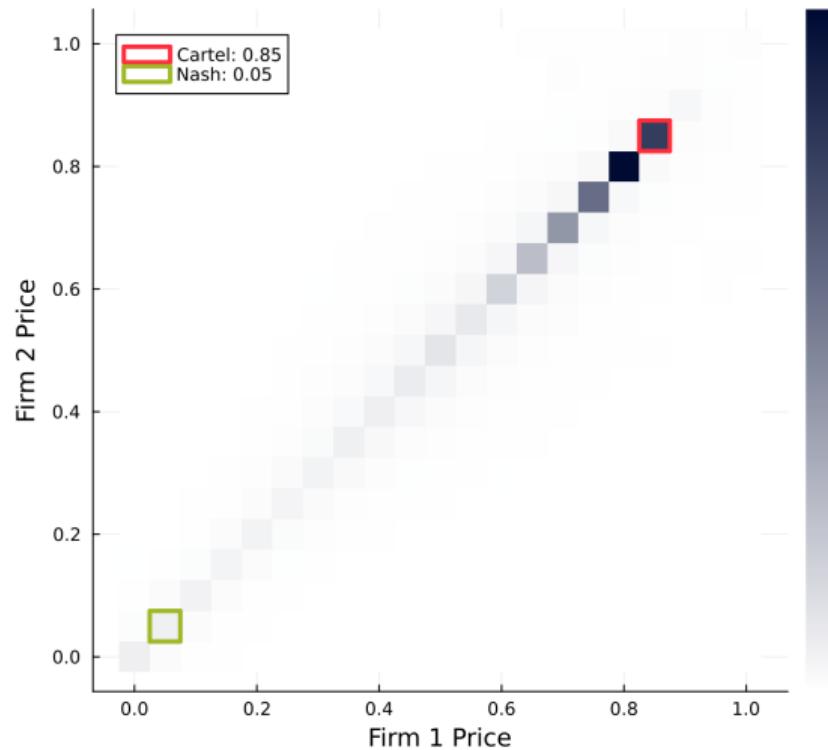
# Results

Robustness - Strategic Environment - Product Differentiation

Product Differentiation:

$$\mu = [0.02 \quad 0.04 \quad 0.06 \quad 0.08 \quad 0.10]$$

- Collusive outcomes robust to degree of product differentiation
- Product homogeneity increases price symmetry
  - Profits more sensitive to being undercut



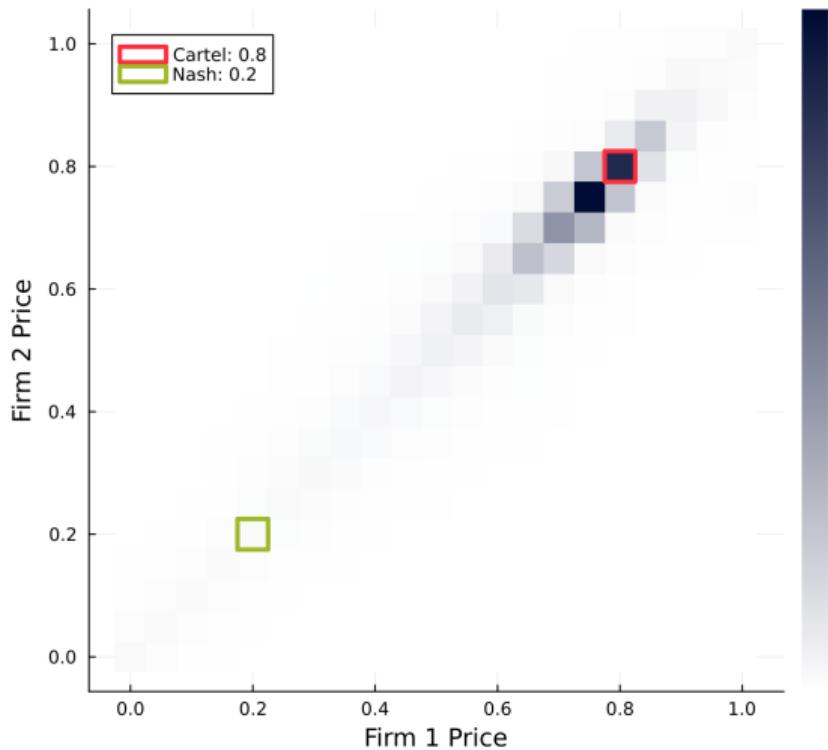
# Results

Robustness - Strategic Environment - Firms

Firms:

$$F = [2 \quad 3 \quad 4]$$

- Difficult to sustain collusion with 3+ firms intermittently deviating
- Consistent with:
  - Theoretical comparative statics
  - Antitrust coordinated effects rule of thumb



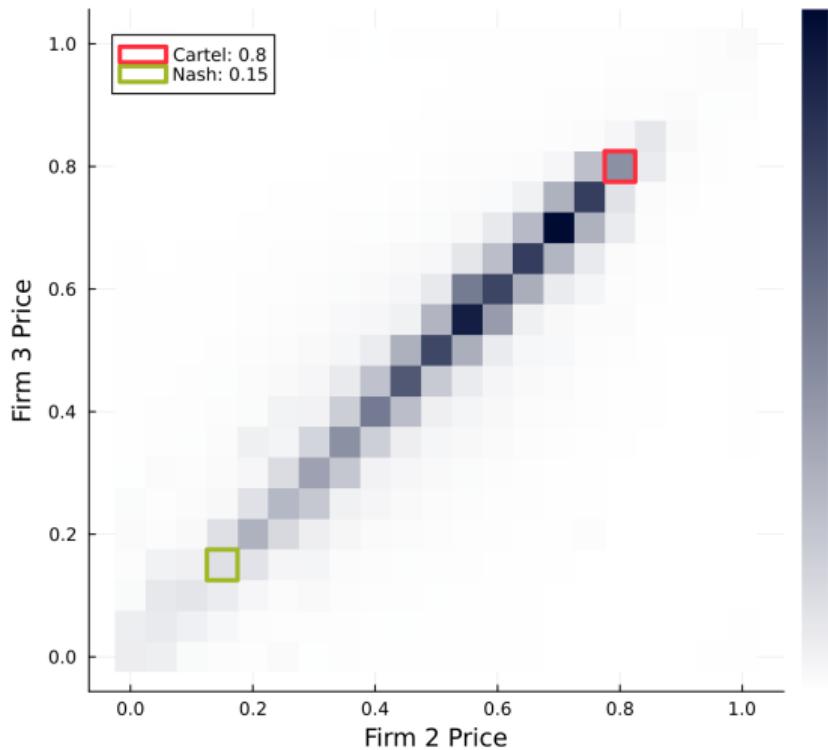
# Results

Robustness - Strategic Environment - Firms

Firms:

$$F = [2 \quad 3 \quad 4]$$

- Difficult to sustain collusion with 3+ firms intermittently deviating
- Consistent with:
  - Theoretical comparative statics
  - Antitrust coordinated effects rule of thumb



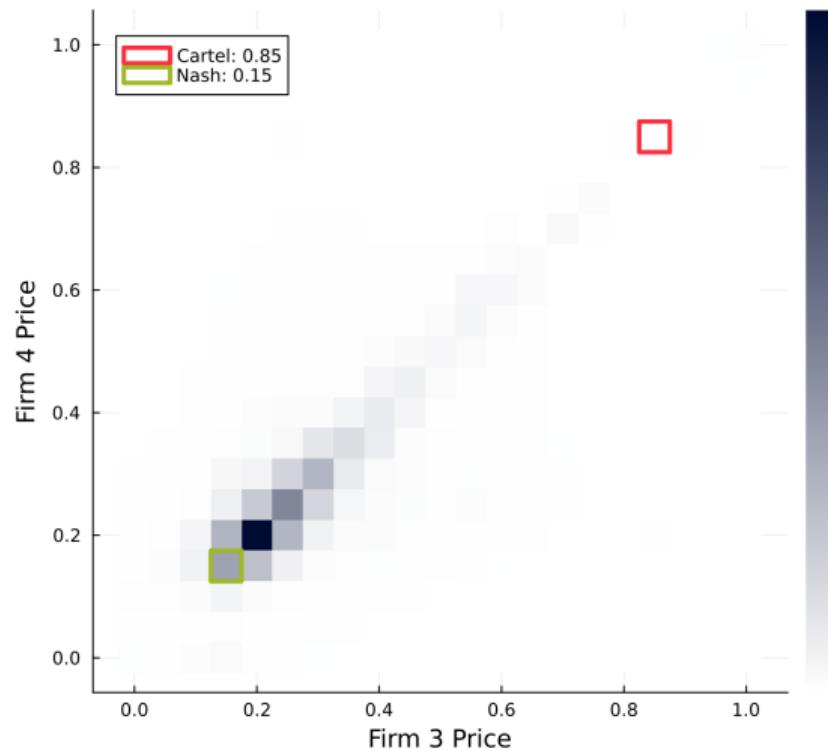
# Results

Robustness - Strategic Environment - Firms

Firms:

$$F = [2 \quad 3 \quad 4]$$

- Difficult to sustain collusion with 3+ firms intermittently deviating
- Consistent with:
  - Theoretical comparative statics
  - Antitrust coordinated effects rule of thumb



# Conclusion

## Summary

- Q-function approximation → efficient learning
  - More perfect collusion
  - Stable collusion after  $\sim 3000$  time periods
- Collusion most sensitive to number of firms
  - Difficult to coordinate with 3+ firms intermittently deviating
- Sensible algorithm parameters required for collusion
  - Sufficient discount factor
  - Sufficient price memory
  - Efficient learning & exploration

